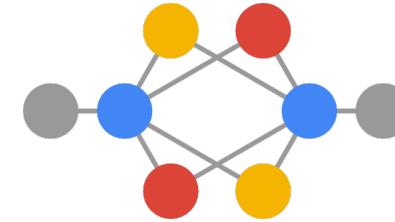
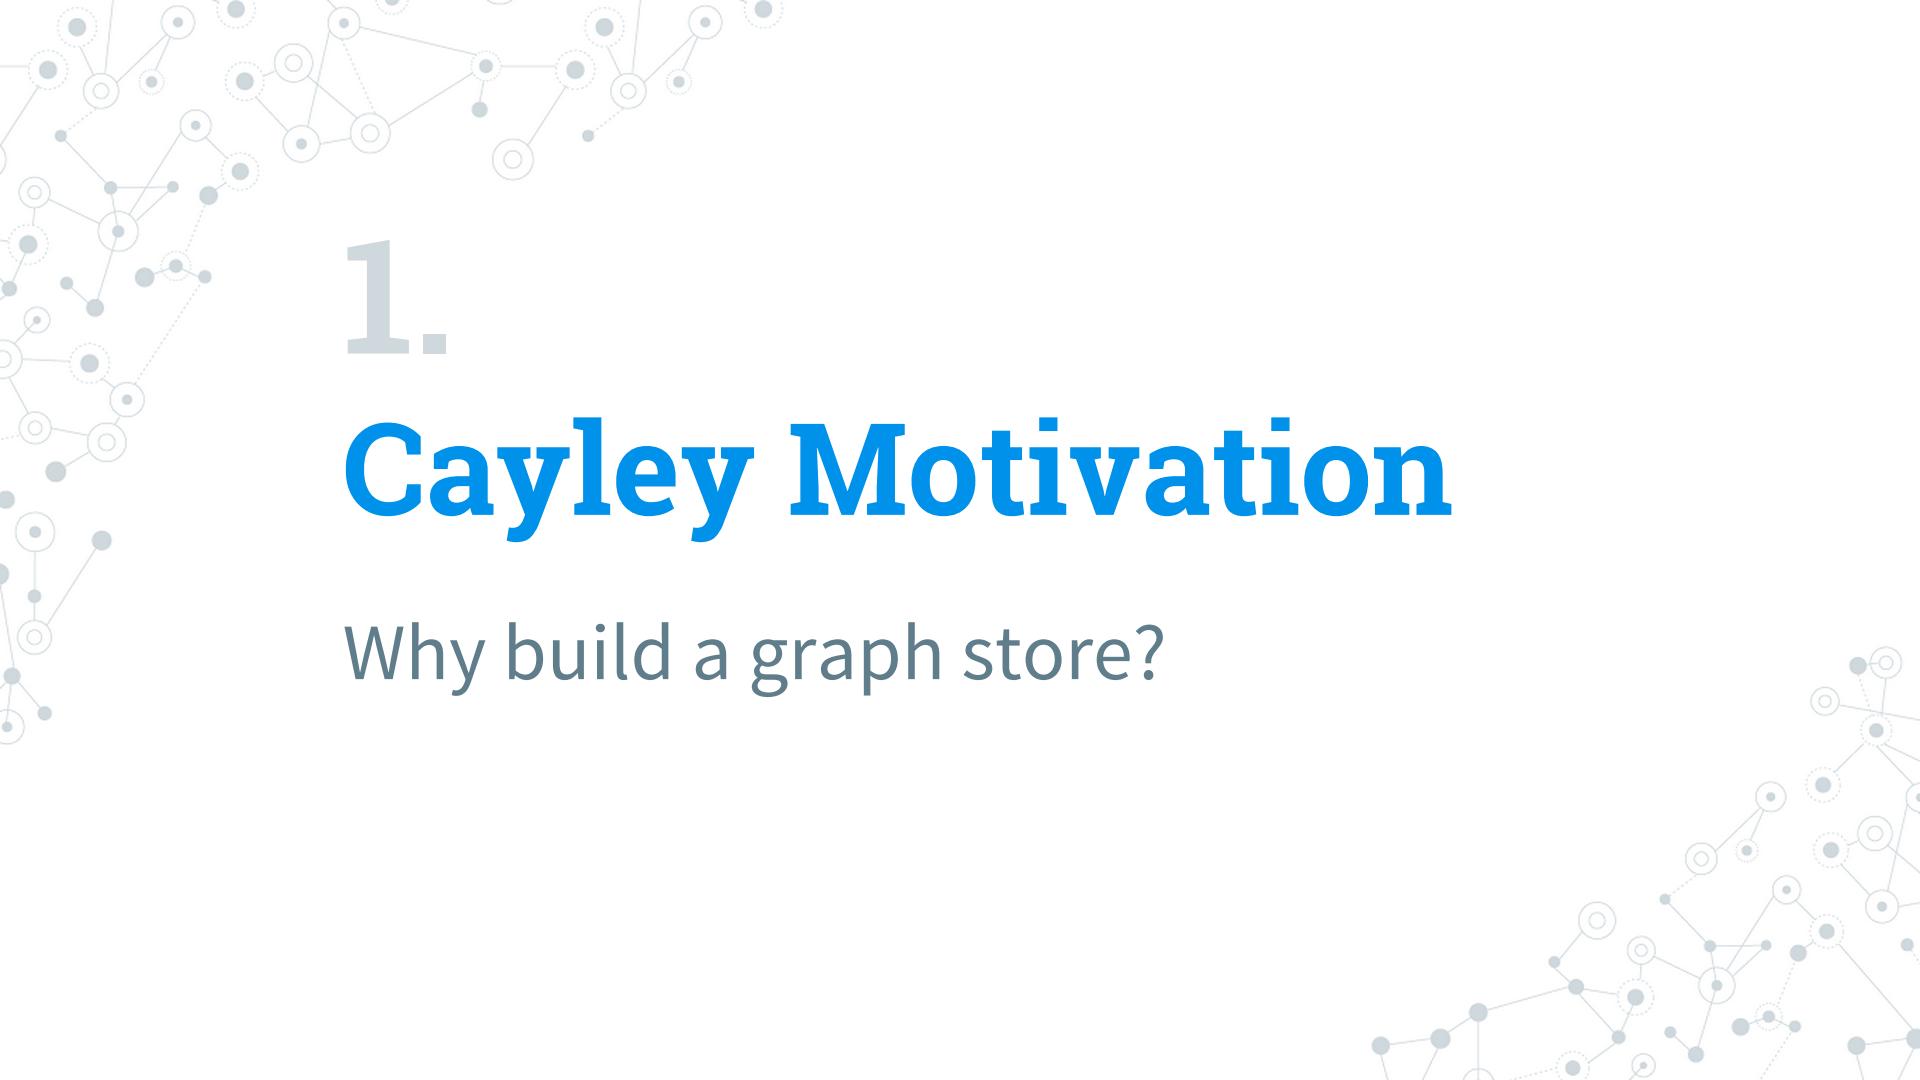


Cayley: Building a Graph Database

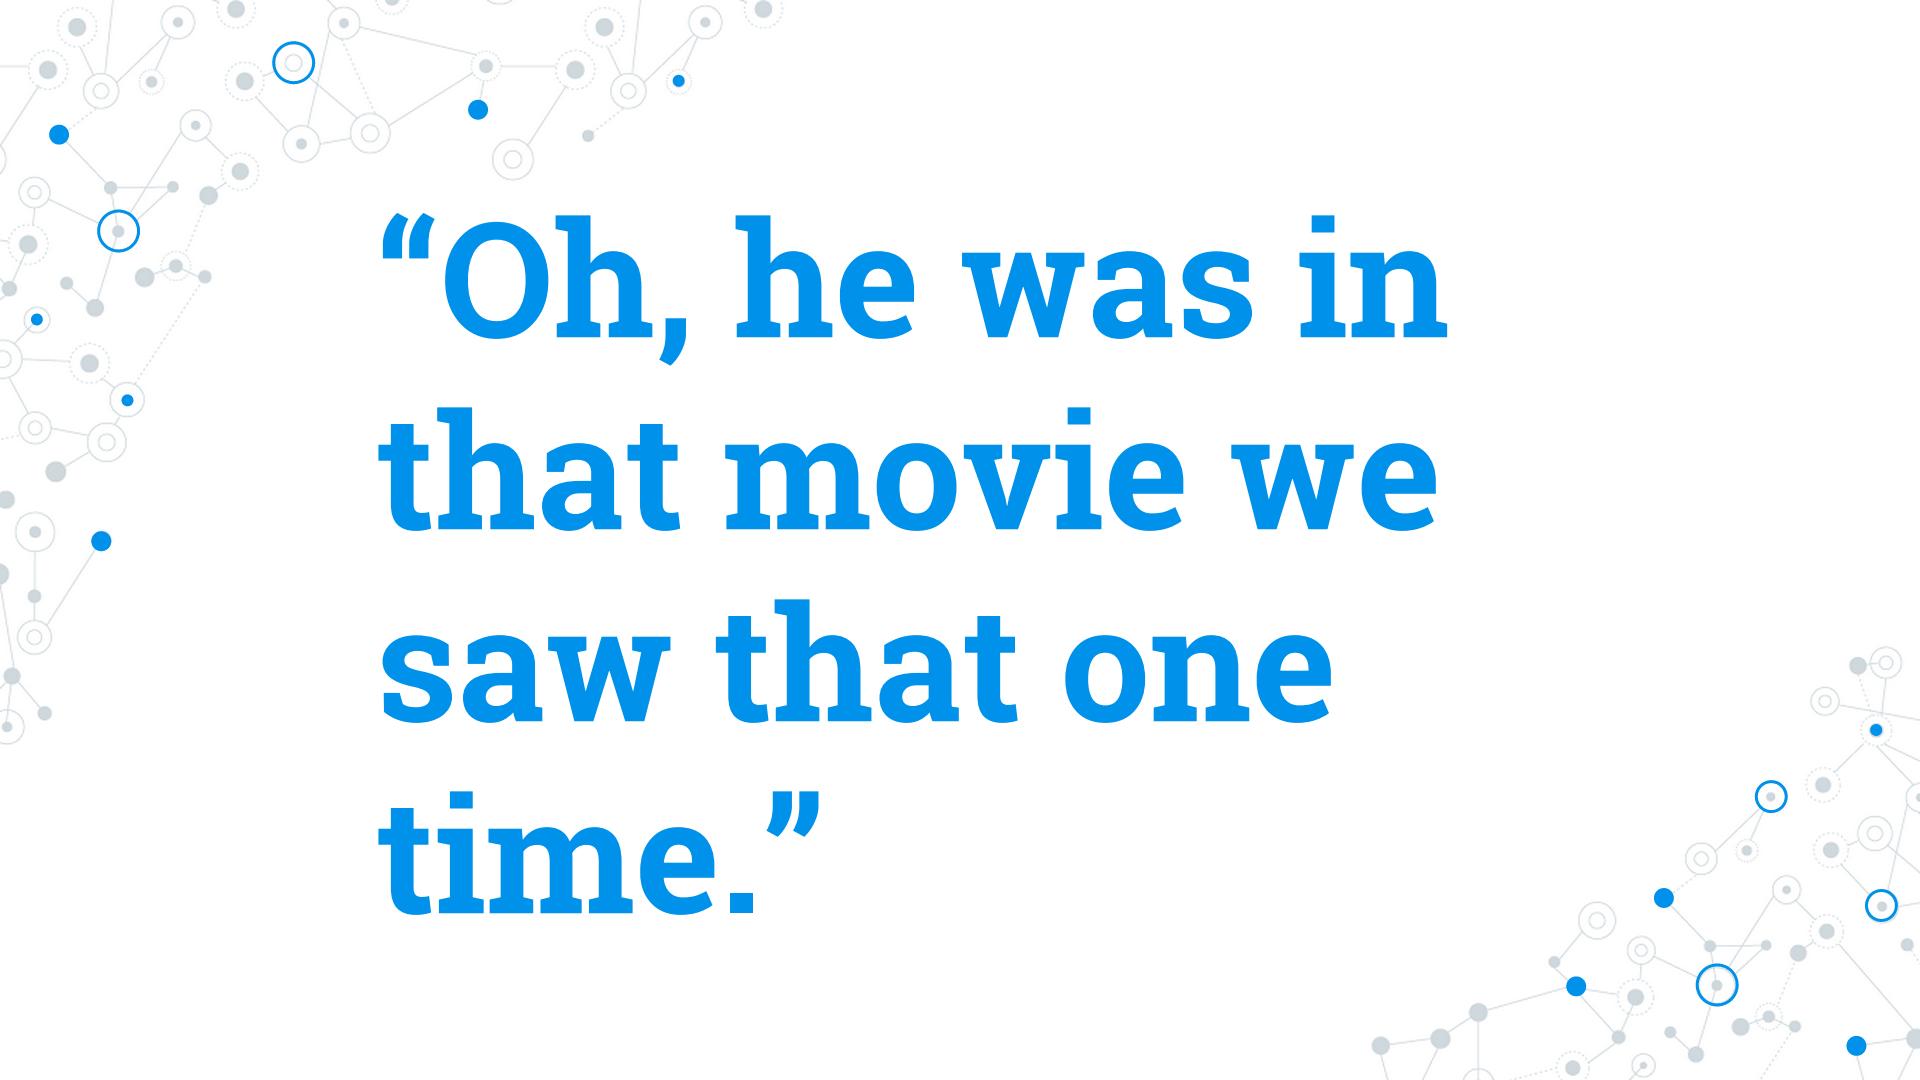




1.

Cayley Motivation

Why build a graph store?



**“Oh, he was in
that movie we
saw that one
time.”**

Movie You're
Watching



Them

You

Event

Some
Other
Movie

We naturally think in terms of connections



Social Networking



Alice

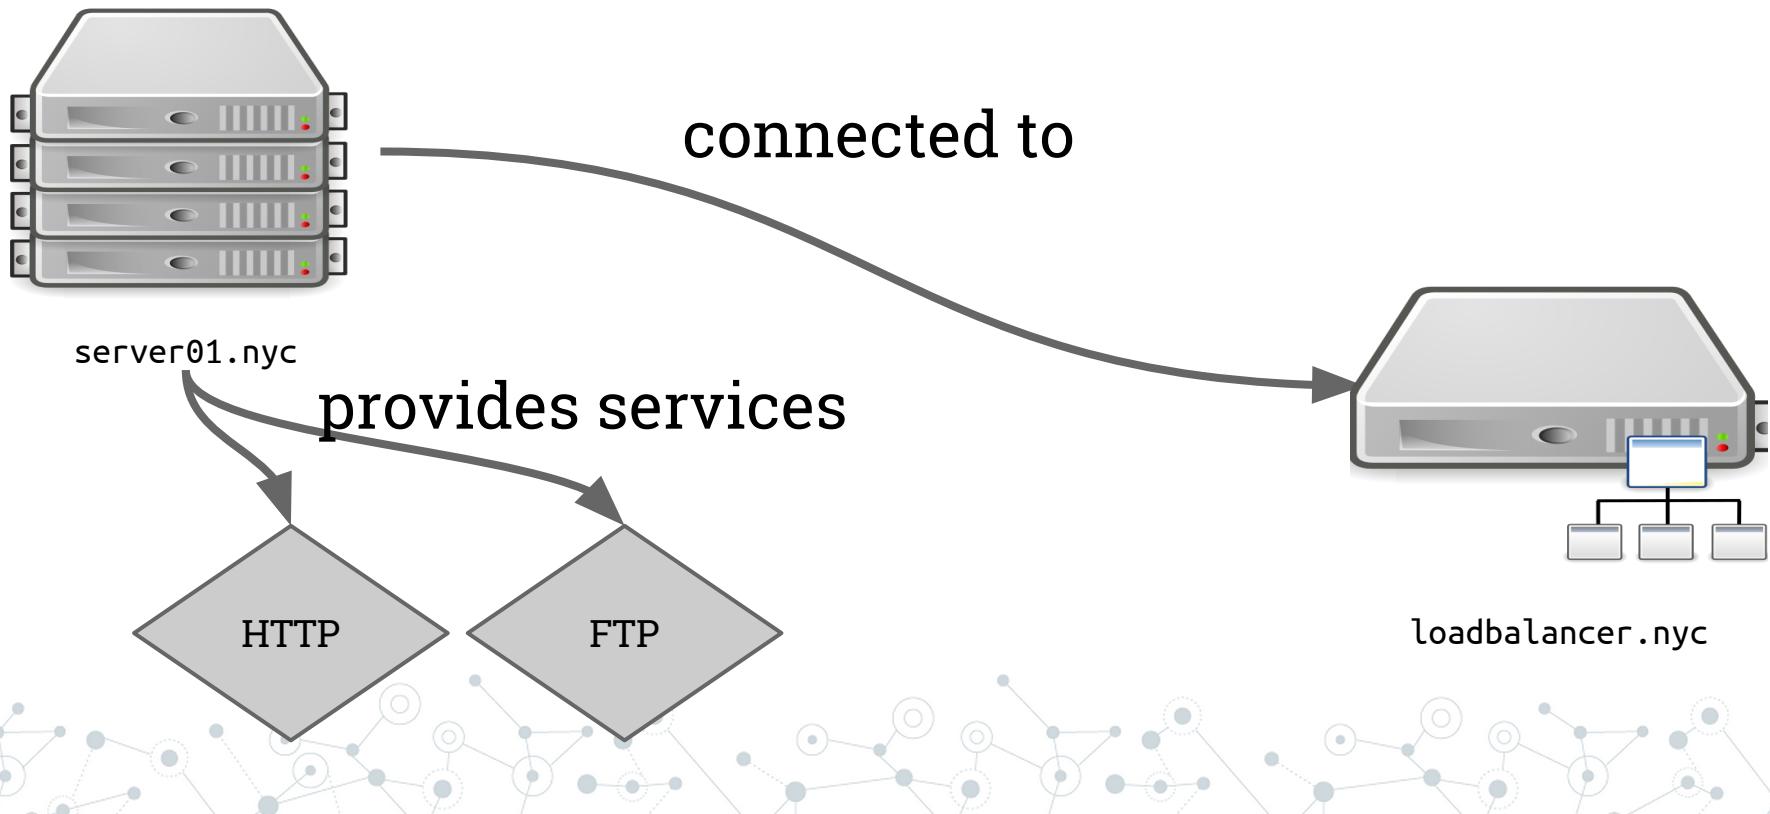
is friends with



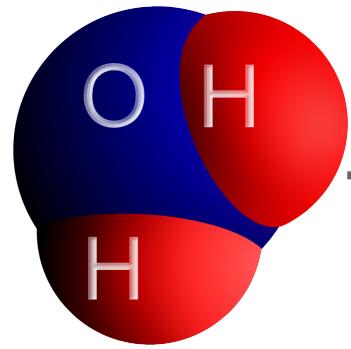
Bob



Computer Networking

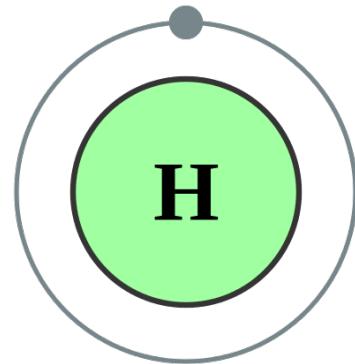


Chemistry



Water

contains element

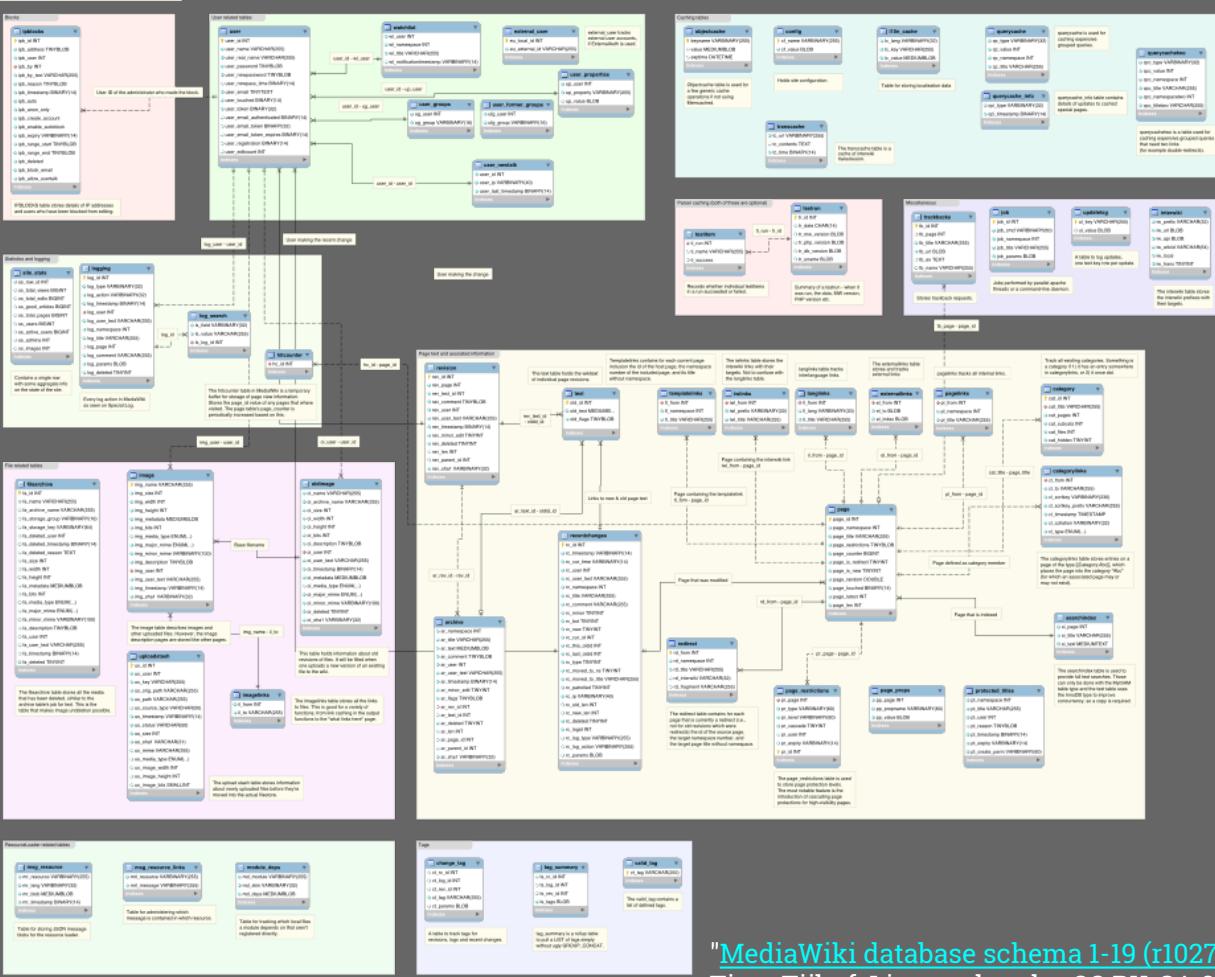


Hydrogen





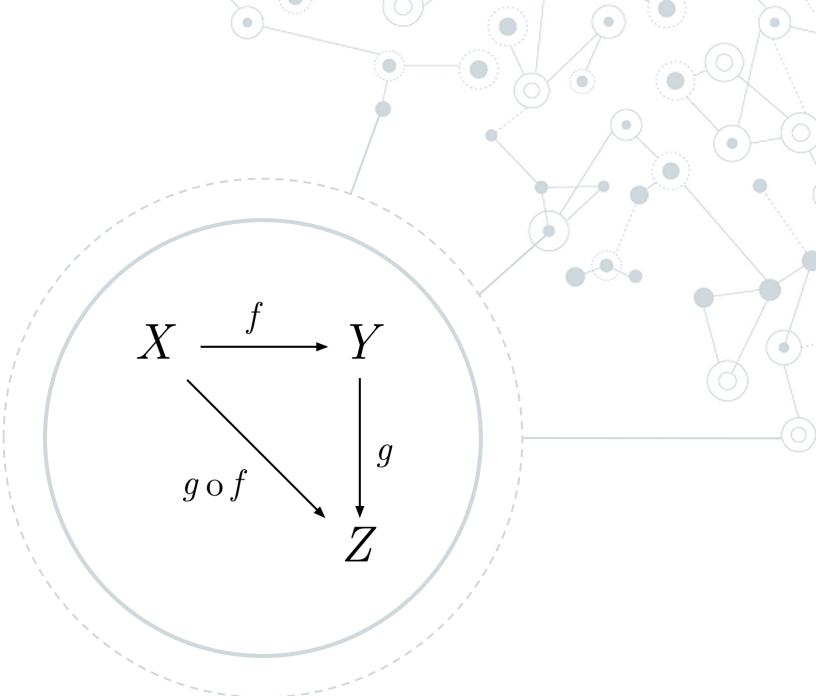
**Today we
construct our
data in tables**



"[MediaWiki database schema 1-19 \(r102798\)](#)" by Timo Tijhof. Licensed under CC BY-SA 3.0 via Wikimedia Commons -

Tables are graphs

Shout out to
Spivak and Kent's Ologs
paper





Microservices: **Why not both?**

**(Can we import a
paradigm?)**



2.

Cayley Origins

A very brief history of Freebase and
Graphd



Freebase was awesome

And graphd was the engine
for the OTG



Giving back to the community

- Freebase shutting down
- Data dumps (and Wikidata)
- Open source some tools...

... but not graphd

graphd was...

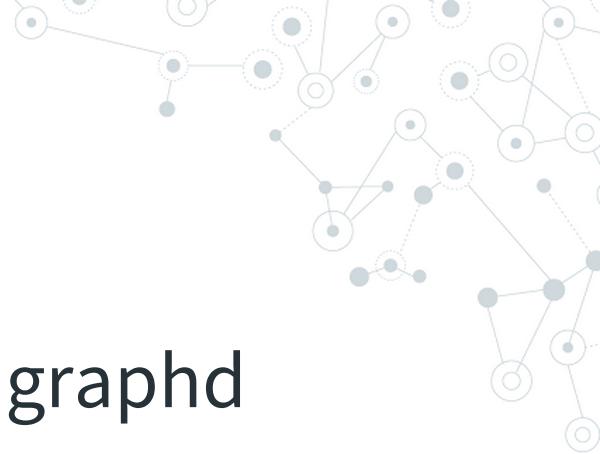
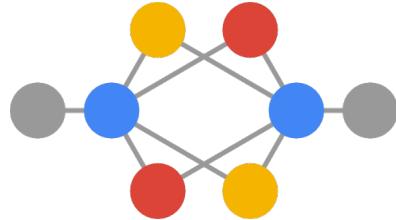
Cool

- Written in C99
- Append-only,
eventually consistent
- Dynamic query
planning
- Custom disk formats
- Replication
- Battle-tested

...but Unapproachable

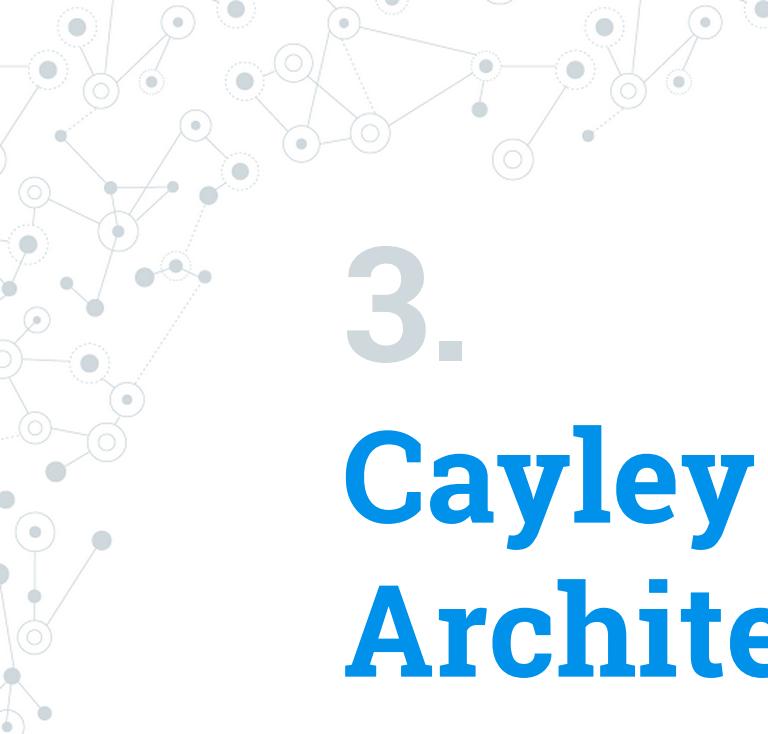
- Written in C99
- Human-unfriendly
query language
- Custom *everything* --
including poll wrapper,
socket sessions,
cooperative
scheduling.

Cayley



- Open-source the novel parts of graphd
- Use Go's amazing standard library
- At least graphd scale (385m primitives pre-Google)
- Promote the graph paradigm
- Agnostic enough for the future
 - Backends
 - Query languages





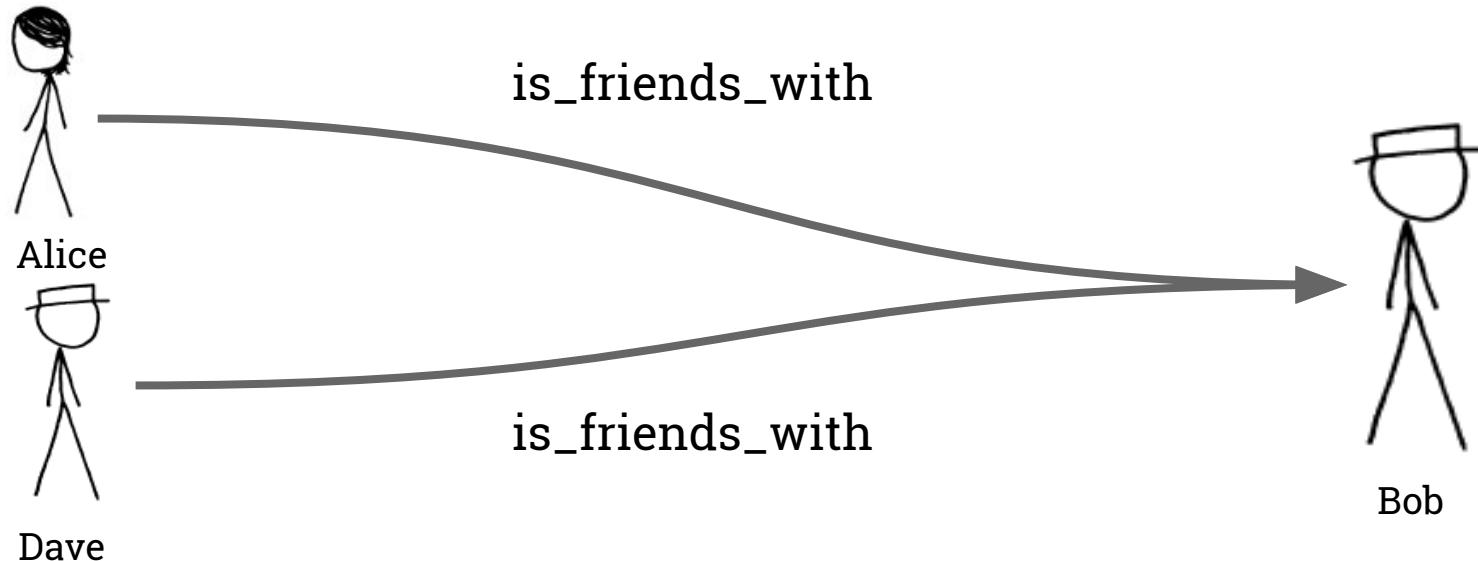
3.

Cayley Architecture

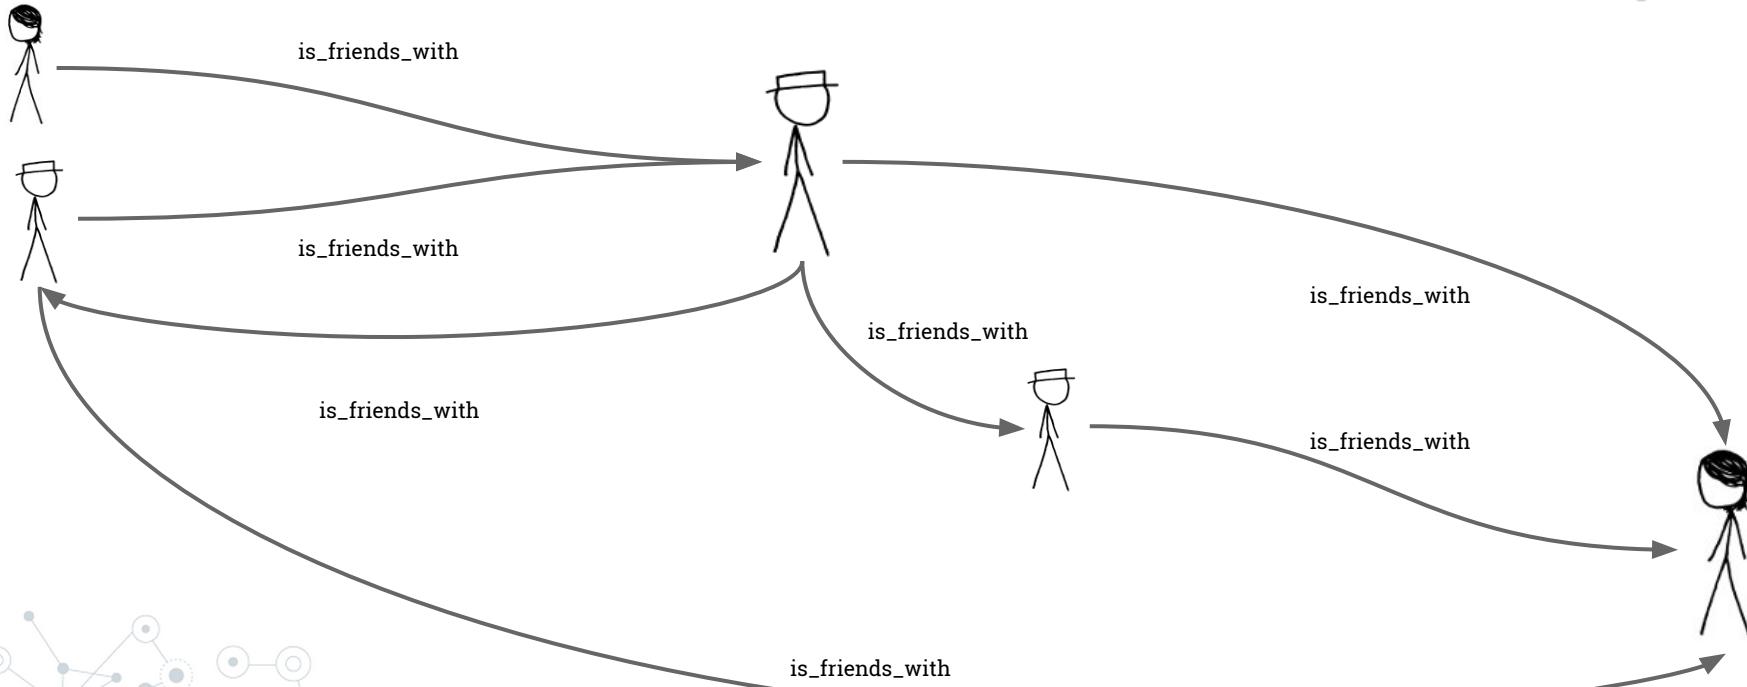
Let's dig into the nuts and bolts



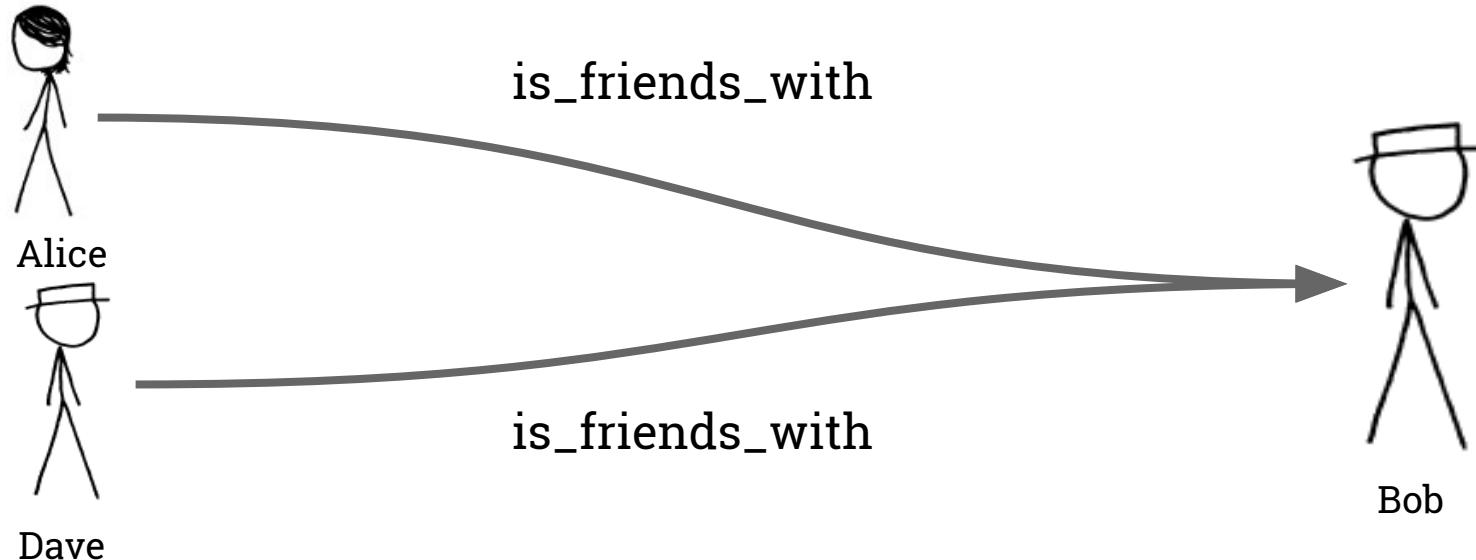
Running Example:



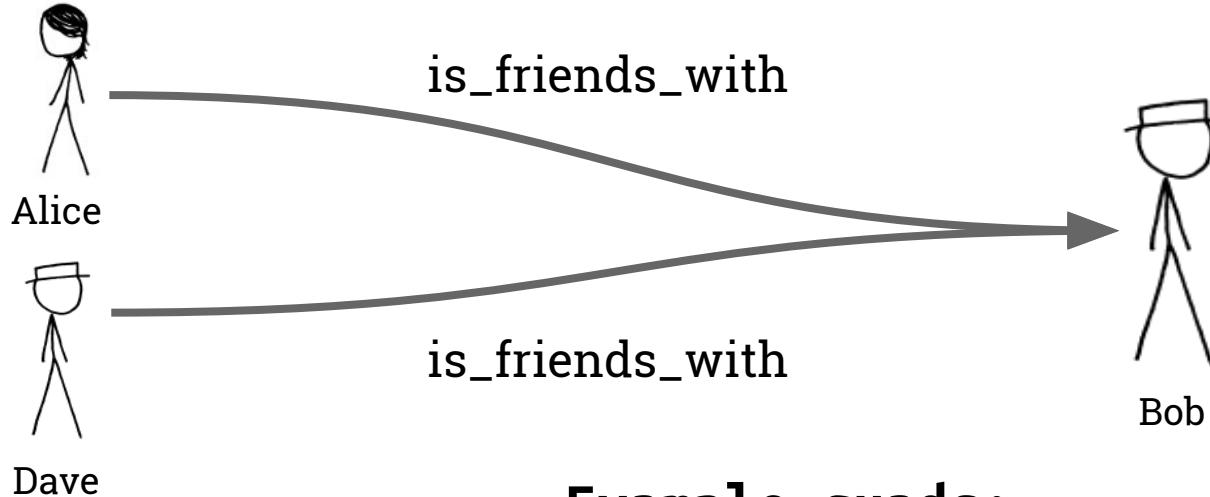
Running Example:



Running Example:



Quads:



Example.quads:

Alice is_friends_with Bob .
Dave is_friends_with Bob .

Quads:

Example.quads:

Alice is_friends_with Bob .

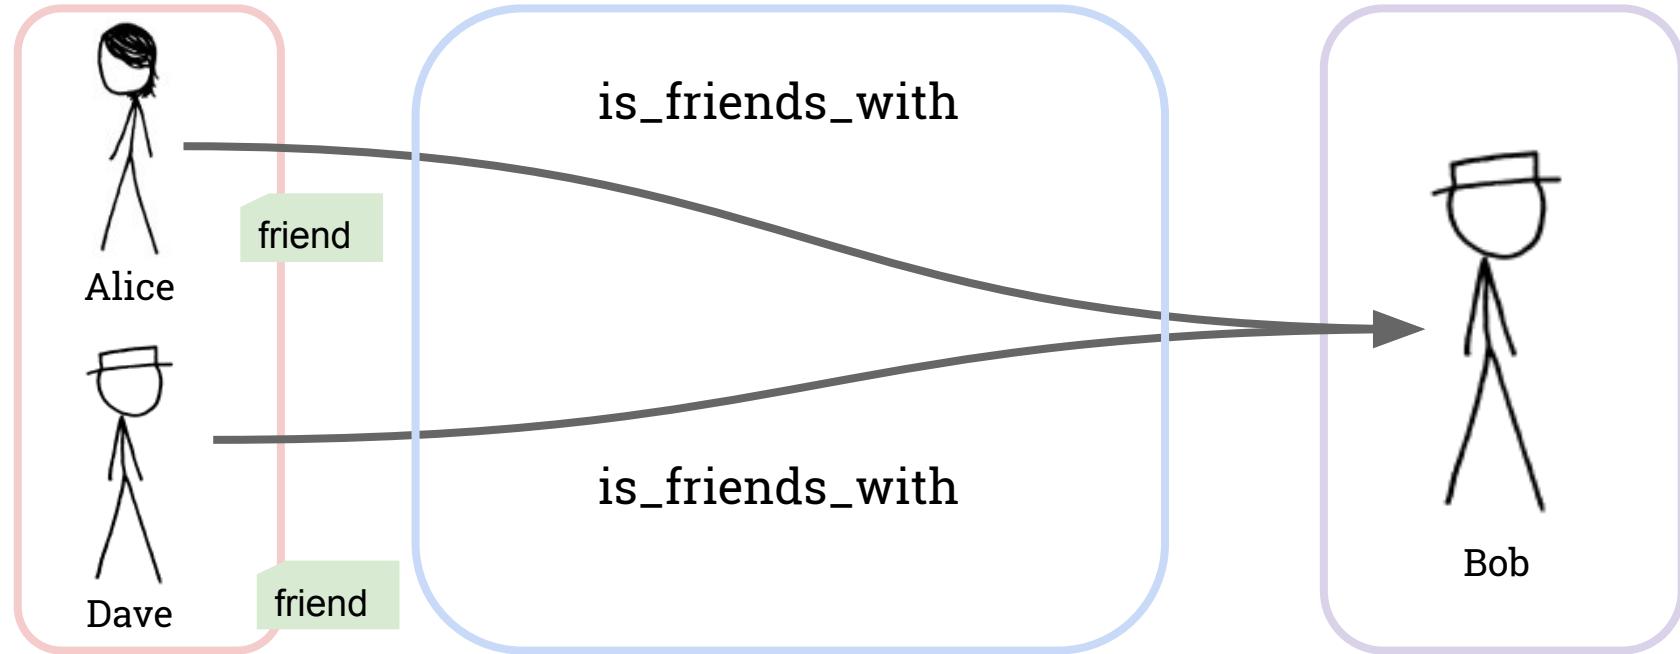
Dave is_friends_with Bob .

Subject

Predicate

Object

Space delimited, fourth field
optional (and ignore it for now)



```
graph.Vertex().Tag("friend").Out("is_friends_with").Is("Bob")
```

Input:

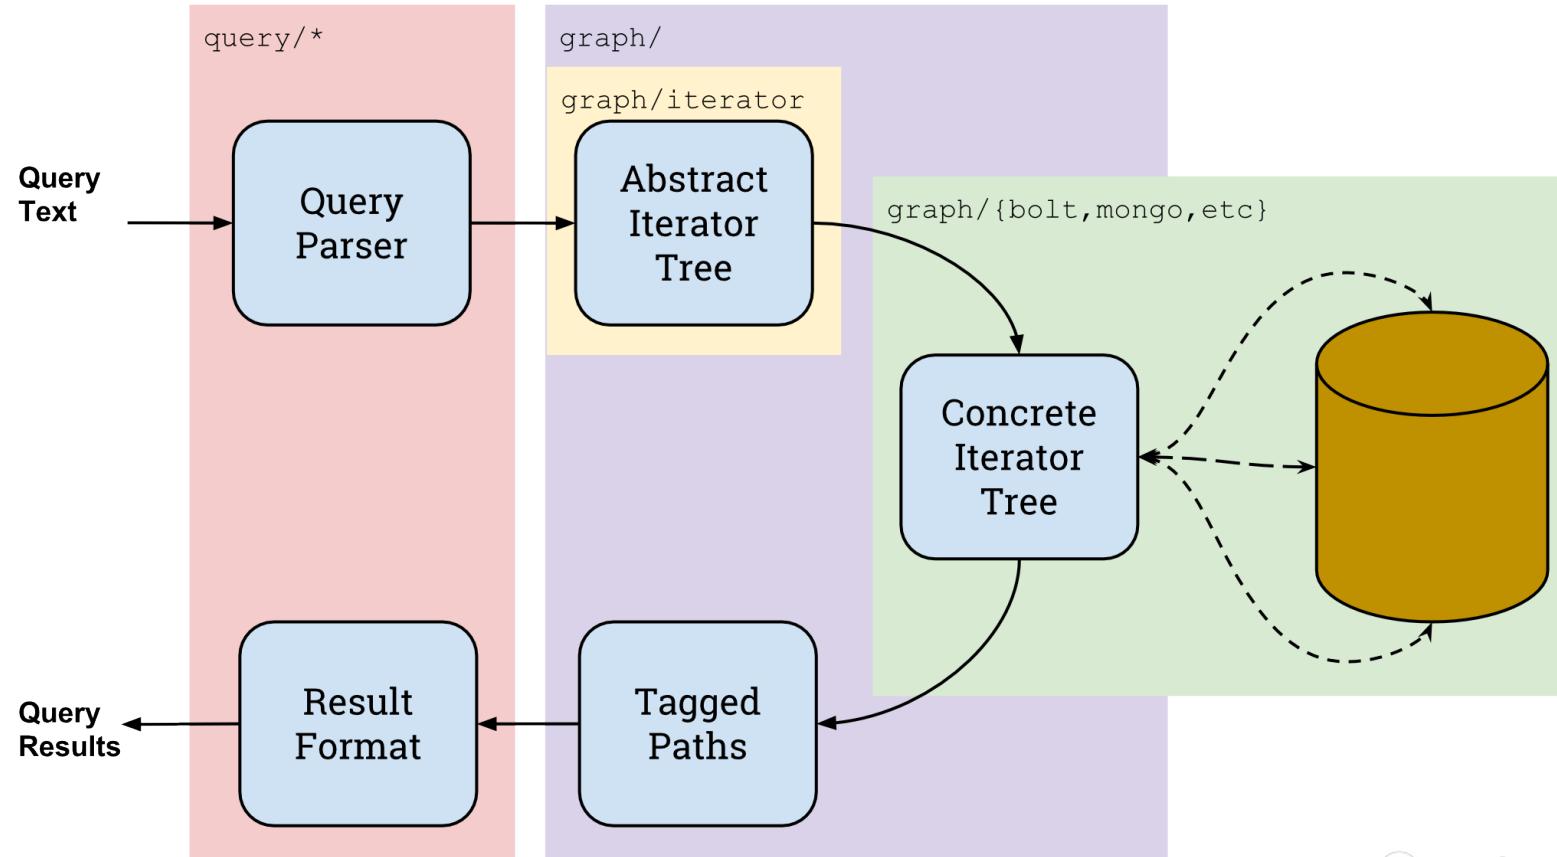
```
graph.Vertex()  
  .Tag("friend")  
  .Out("is_friends_with")  
  .Is("Bob")  
  .All()
```

Output:

```
{  
  "result": [  
    {  
      "friend": "Alice",  
      "id": "Bob"  
    },  
    {  
      "friend": "Dave",  
      "id": "Bob"  
    }  
  ]  
}
```



Life of a Graph Query





3a.

Query Languages

query/*



Role: Parser

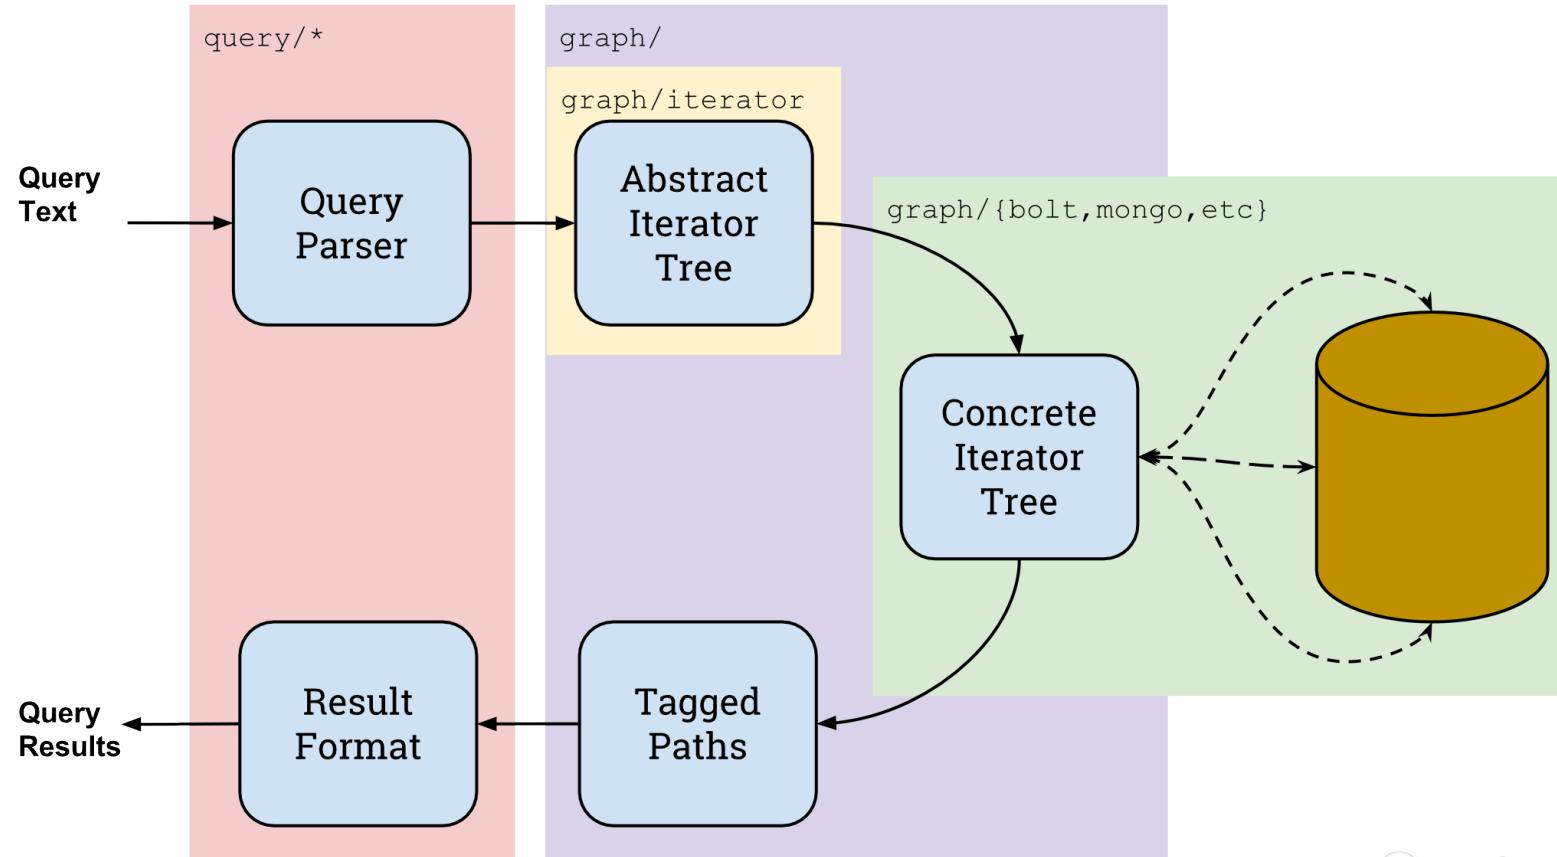
- Create a user session
- Build an abstract iterator tree based on user input.
- Evaluate this tree and return results

Primary Interface:

query/session.go:

```
// Does this even parse, really? Or do I need to keep reading?  
Parse(query string) (ParseResult, error)  
  
// Parse and evaluate, giving me paths back on the channel  
Execute(query string, res chan interface{}, limit int)  
  
// Stringify said results for the target audience  
Format(interface{}) string
```

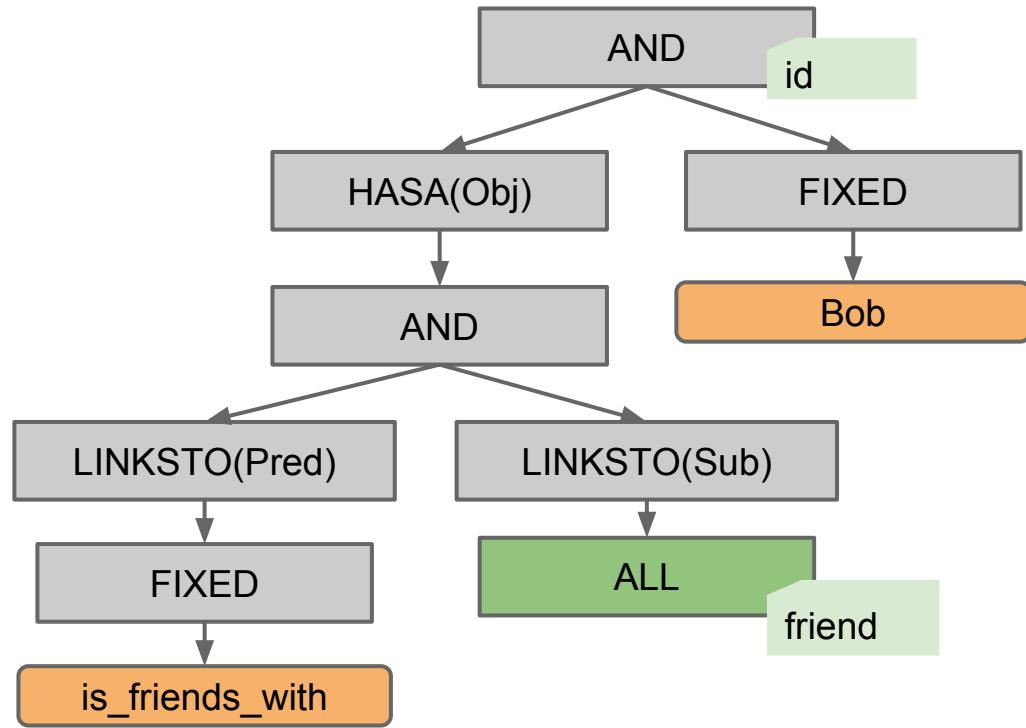
Life of a Graph Query



Input:

```
graph.Vertex()  
.Tag("friend")  
.Out("is_friends_with")  
.Is("Bob")  
.All()
```

Output:





3b.

Abstract Iterators

graph/iterator

Role: Represent Sets

- Placeholder for unevaluated sets of either nodes or quads (links)
- Precise set and map operations
- Be backend-independent (“abstract”)
- Hold tags

An iterator glossary



AND

- Intersection: “*In this set AND in this set*”



OR

- Union: “*In this set OR in this set*”



FIXED

- An explicit set in memory, often size 1: “*Exactly ‘A’*”



ALL

- All nodes in the database: “*All of it!*”

Iterator traversal duals



LINKSTO

- Set of Nodes -> Set of Links



HASA

- Set of Links -> Set of Nodes



An iterator glossary



LINKSTO(direction)

- All quads that have a set of nodes in the direction field
- “*All LINKS TO Alice where Alice is the subject*”



HASA(direction)

- All nodes given by the direction field of a set of quads.

“*This link HAS A object, which is Bob*”

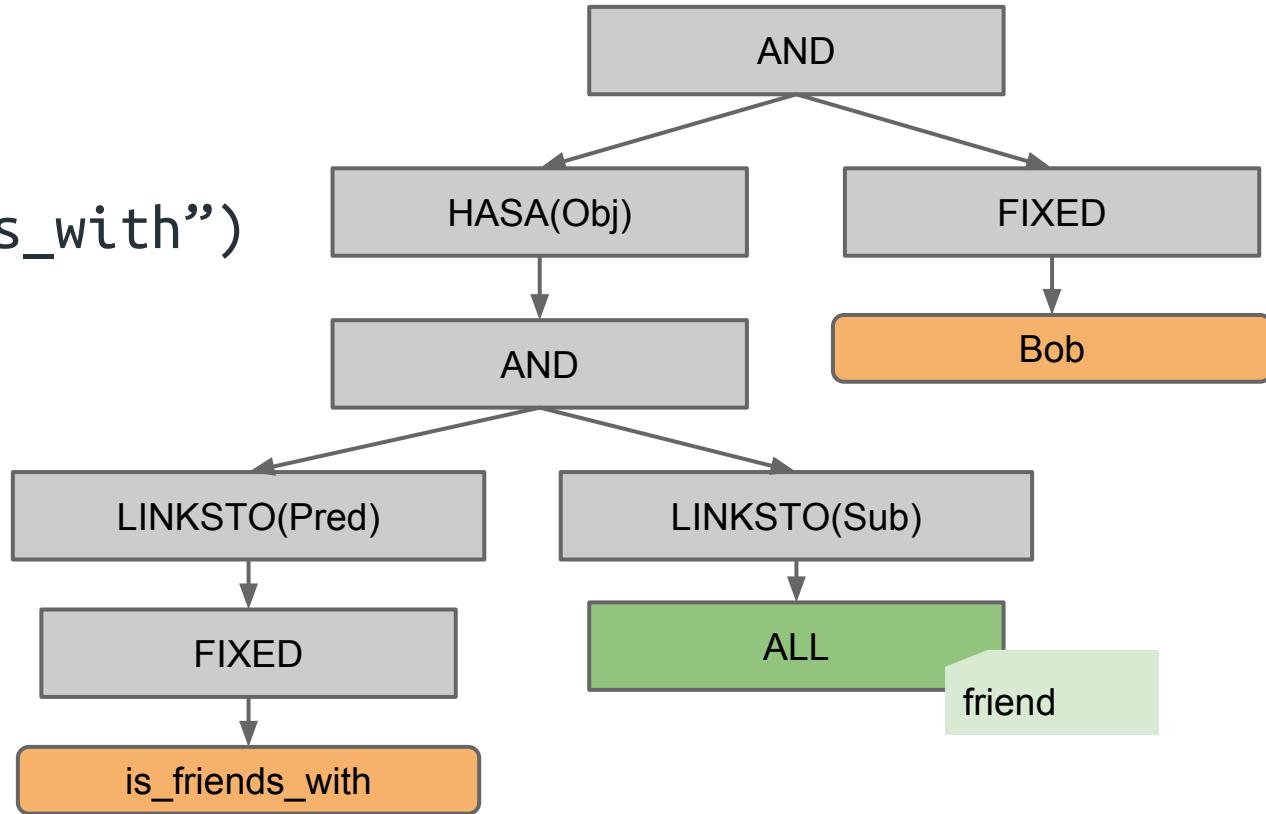


```
graph.Vertex()
```

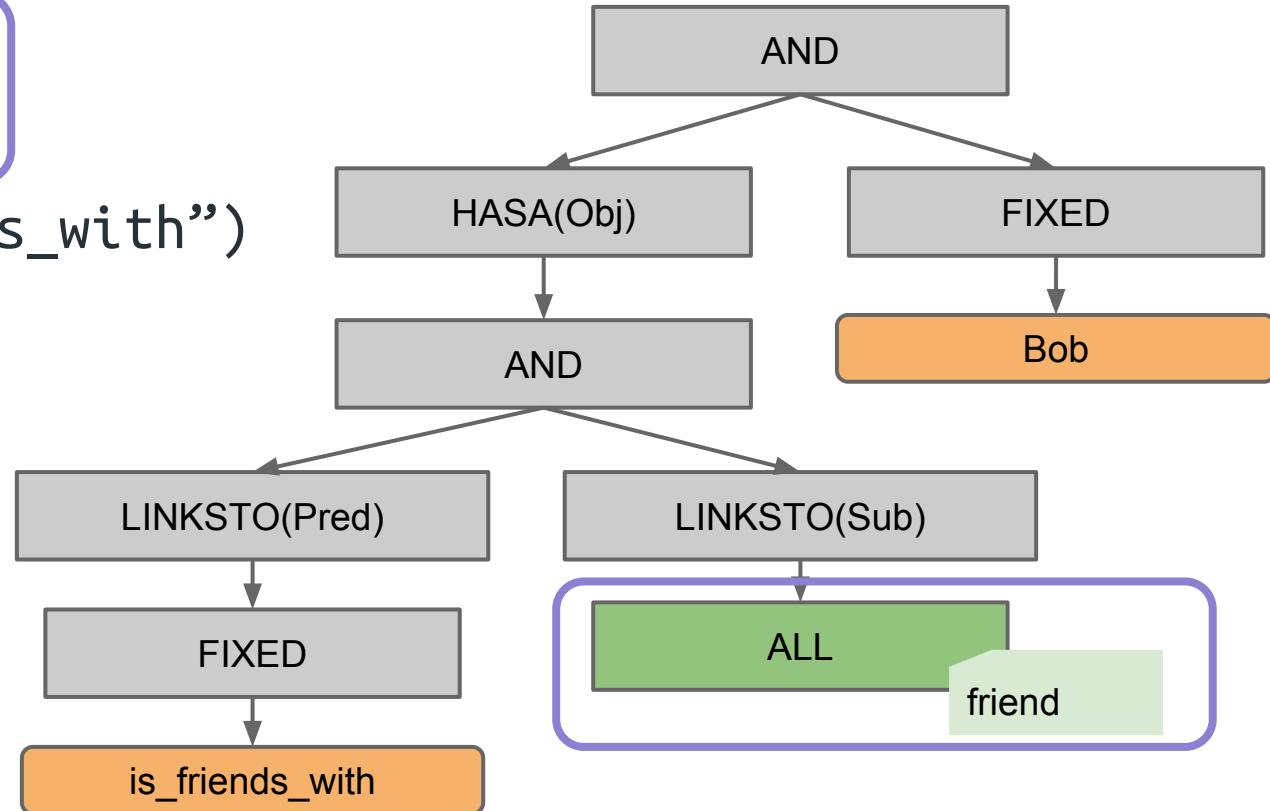
```
.Tag("friend")
```

```
.Out("is_friends_with")
```

```
.Is("Bob")
```



```
graph.Vertex()  
  .Tag("friend")  
  .Out("is_friends_with")  
  .Is("Bob")
```



```
graph.Vertex()
```

```
.Tag("friend")
```

```
.Out("is_friends_with")
```

```
.Is("Bob")
```



```
HASA(Obj)
```

```
FIXED
```

```
AND
```

```
Bob
```

```
LINKSTO(Pred)
```

```
LINKSTO(Sub)
```

```
FIXED
```

```
ALL
```

```
is_friends_with
```

```
friend
```

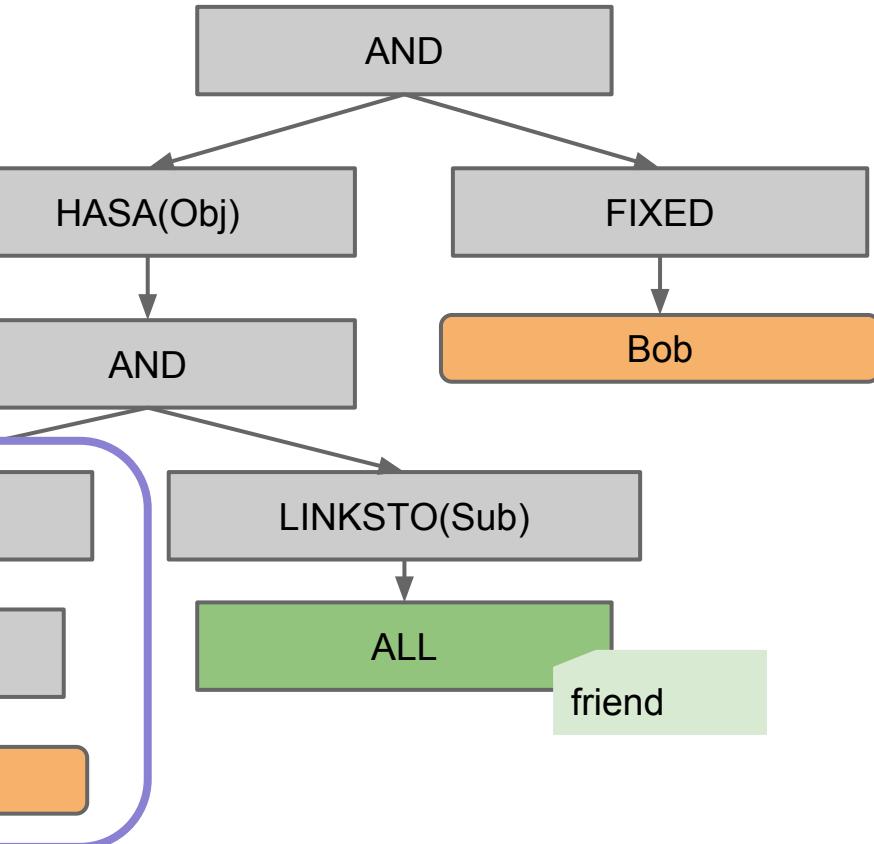


```
graph.Vertex()
```

```
.Tag("friend")
```

```
.Out("is_friends_with")
```

```
.Is("Bob")
```



```
graph.Vertex()
```

```
.Tag("friend")
```

```
.Out("is_friends_with")
```

```
.Is("Bob")
```



```
HASA(Obj)
```

```
FIXED
```

```
AND
```

```
Bob
```

```
LINKSTO(Pred)
```

```
LINKSTO(Sub)
```

```
FIXED
```

```
ALL
```

```
is_friends_with
```

```
friend
```

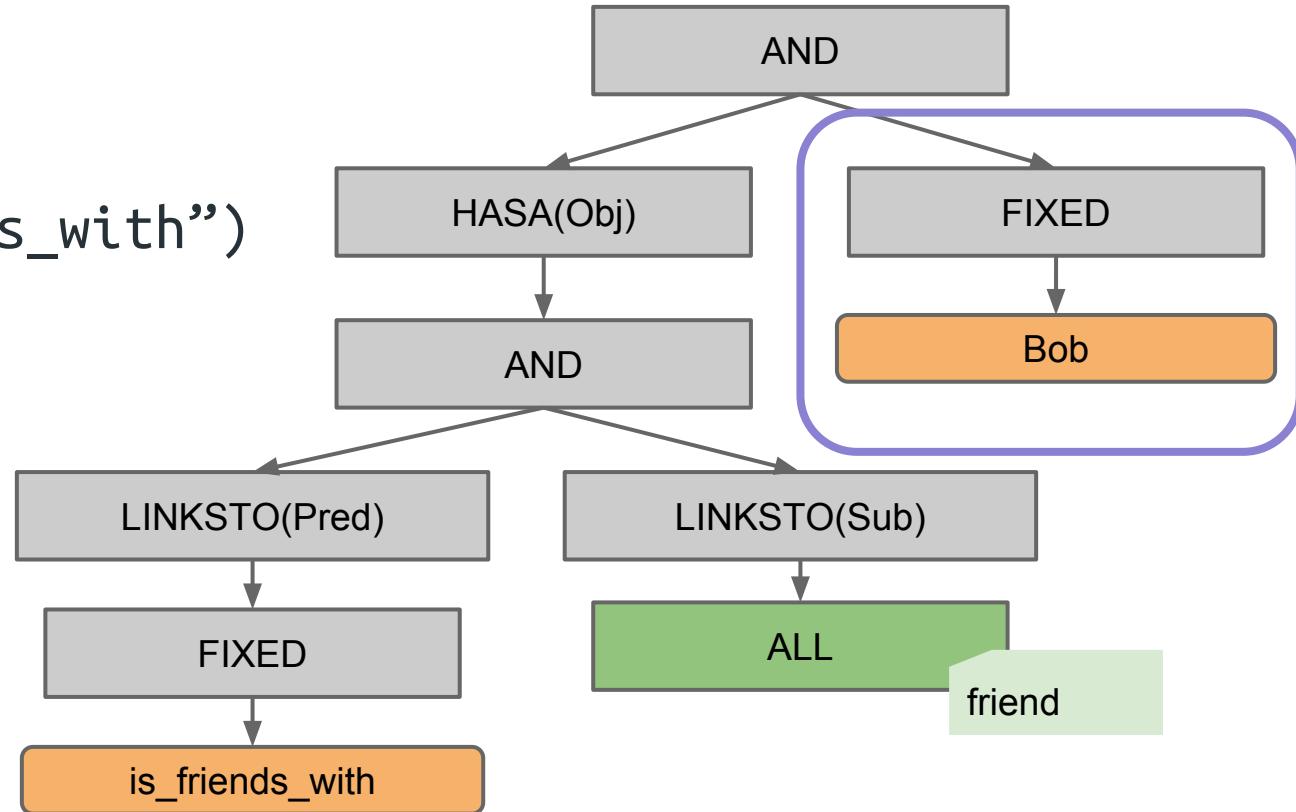


```
graph.Vertex()
```

```
.Tag("friend")
```

```
.Out("is_friends_with")
```

```
.Is("Bob")
```



Using Iterators:

graph/iterator.go:

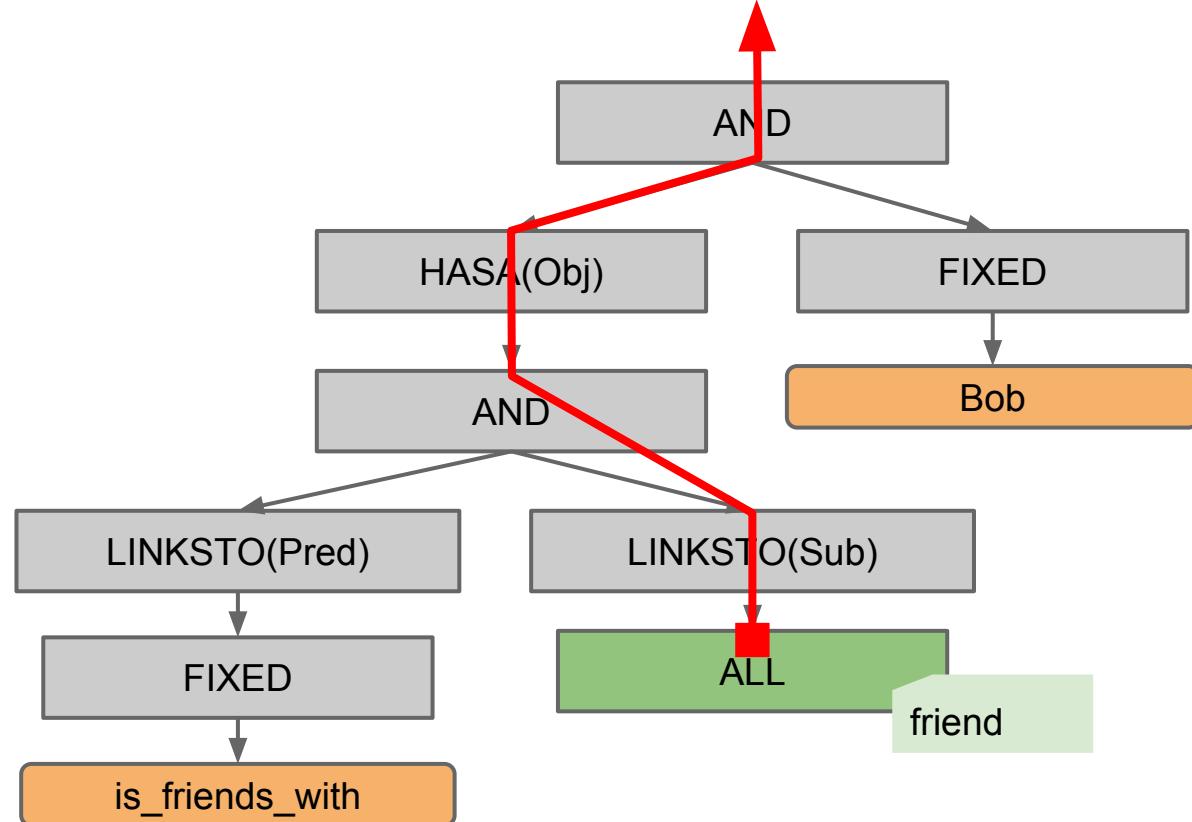
```
// Get the Value token at the current iterator position
Result() Value

// Move the iterator to the next result value
Next() bool

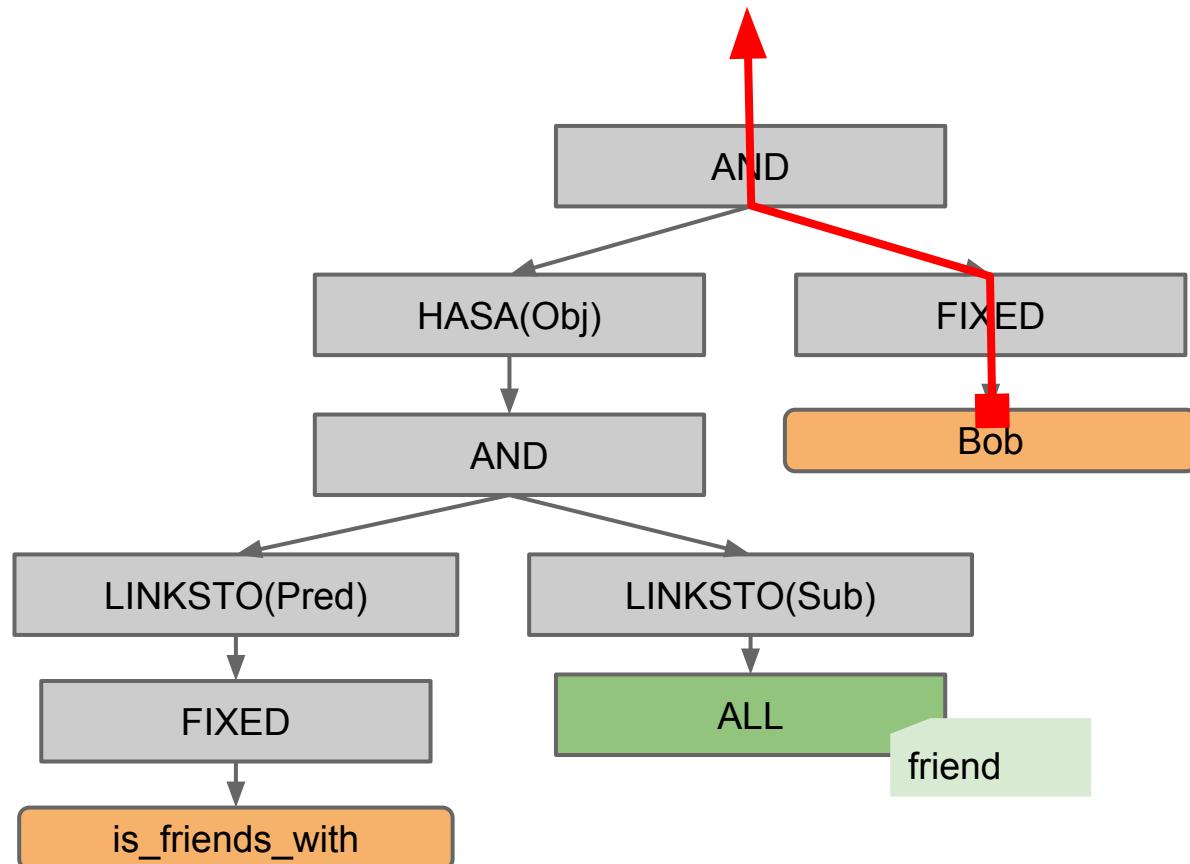
// Is there some other path in the tree that would
// give us the same Result()? If so, advance the subiterators.
NextPath() bool

// Does the set this iterator represents contain this token?
Contains(Value) bool
```

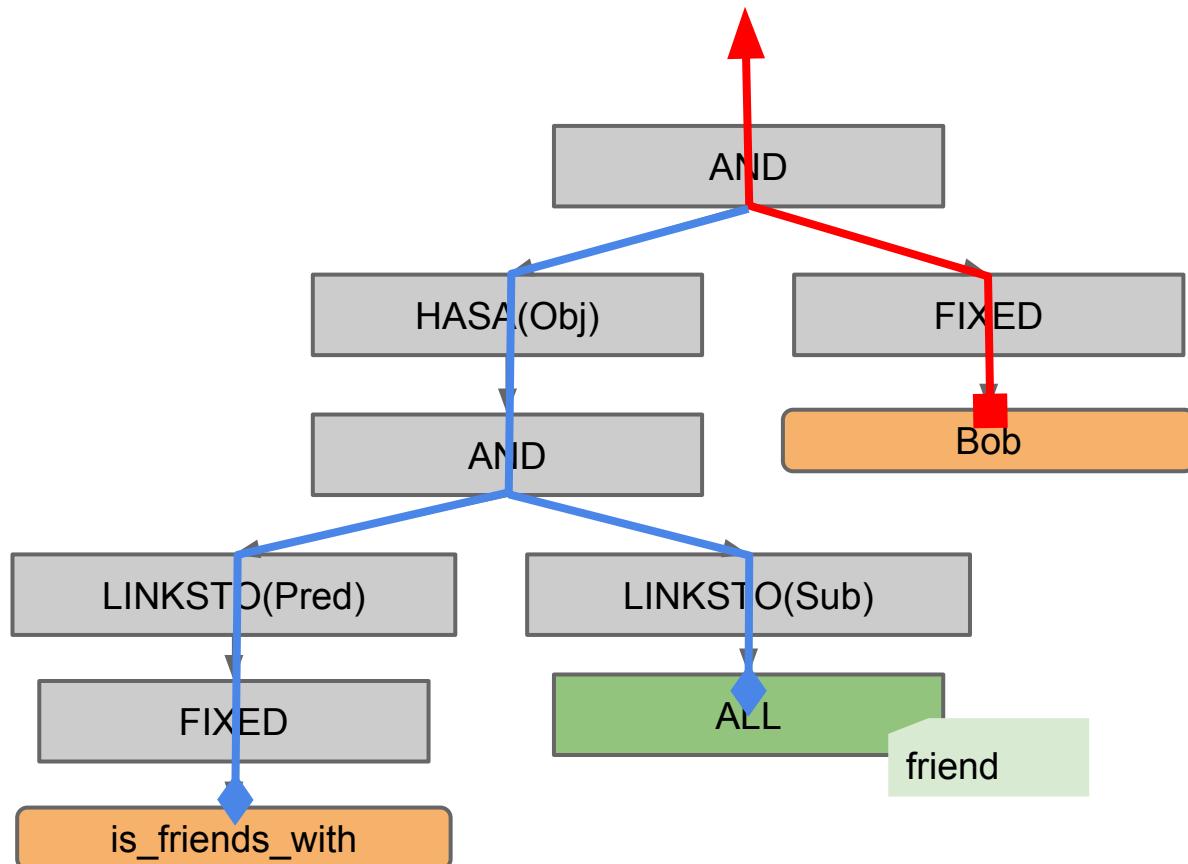
Avoid the naive path



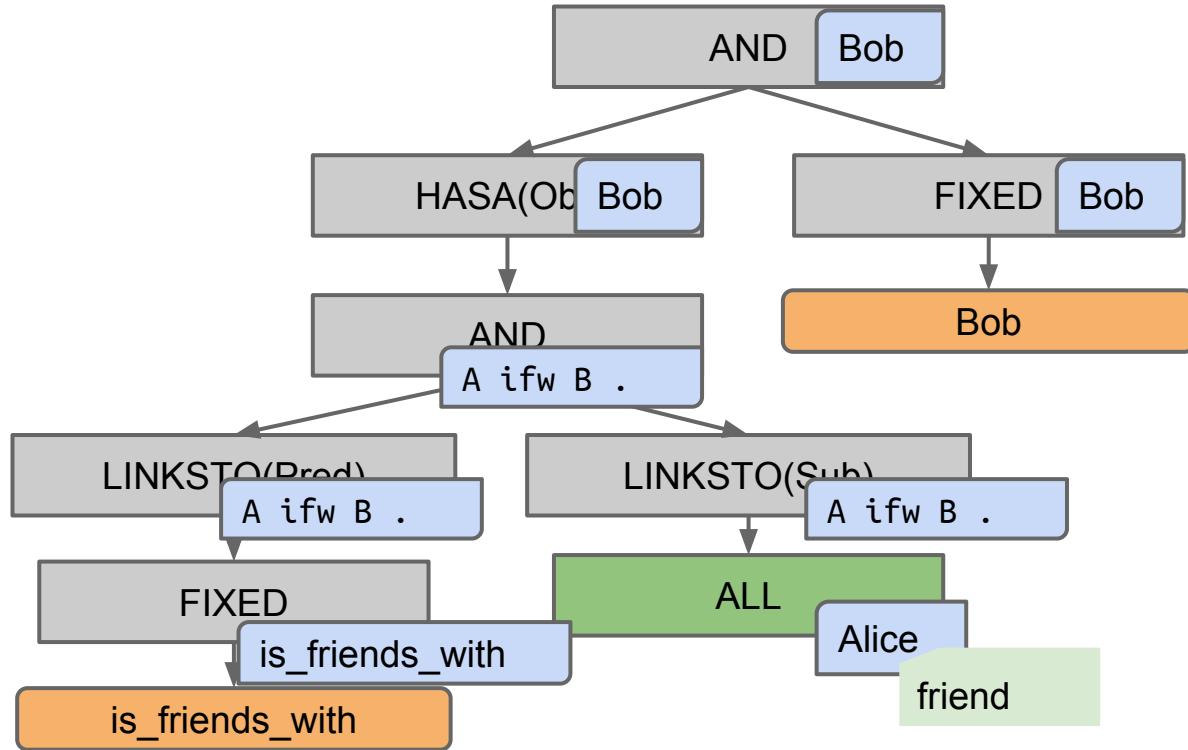
Next()
from the
easy
part...



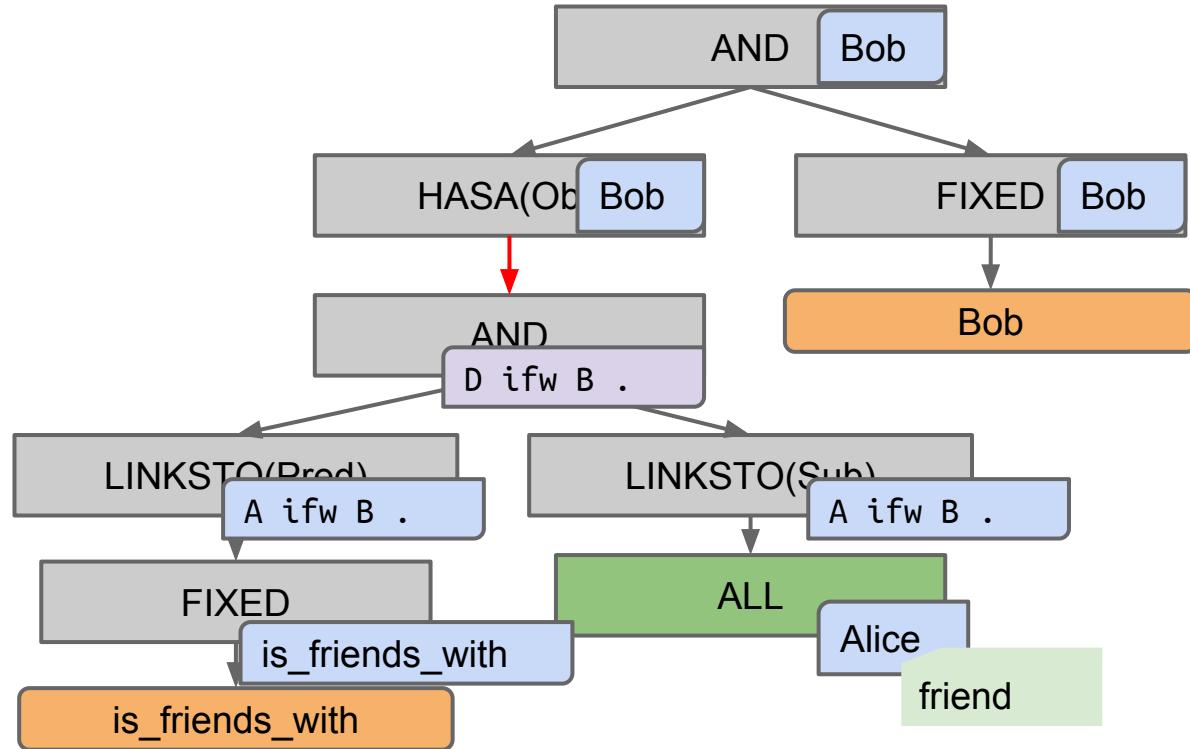
Contains() checks the rest



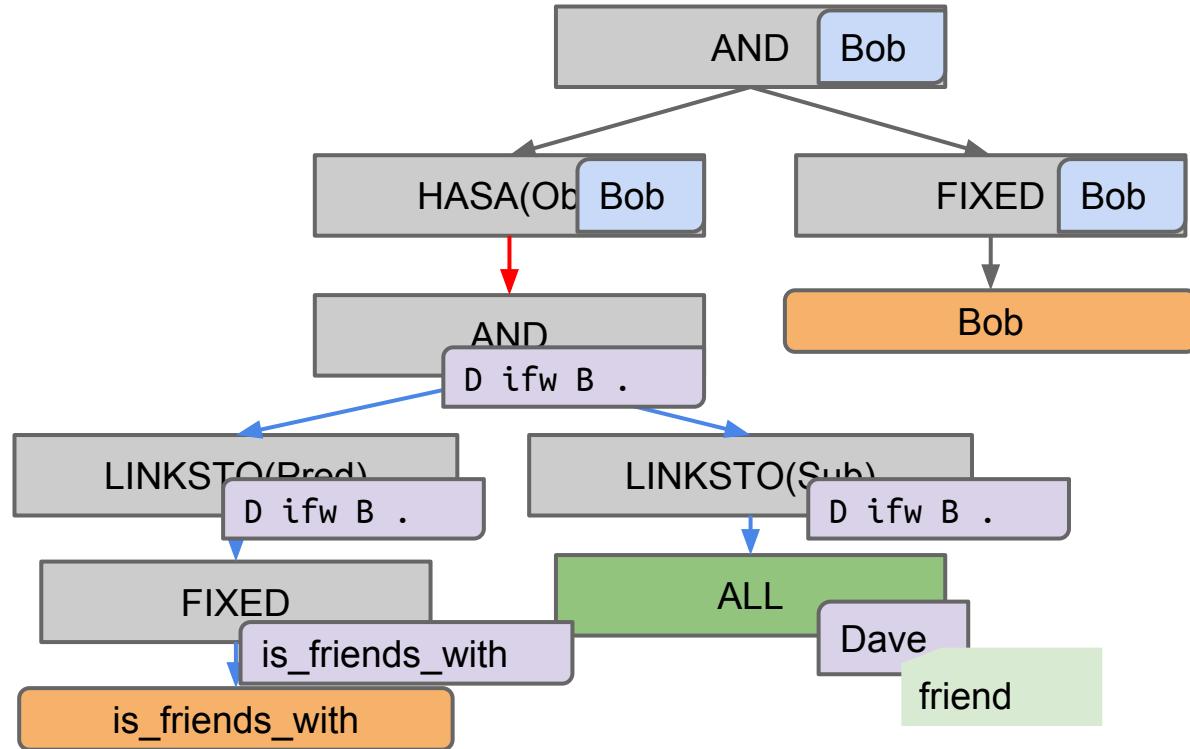
Result()



NextPath() to find the other friends



NextPath() to find the other friends



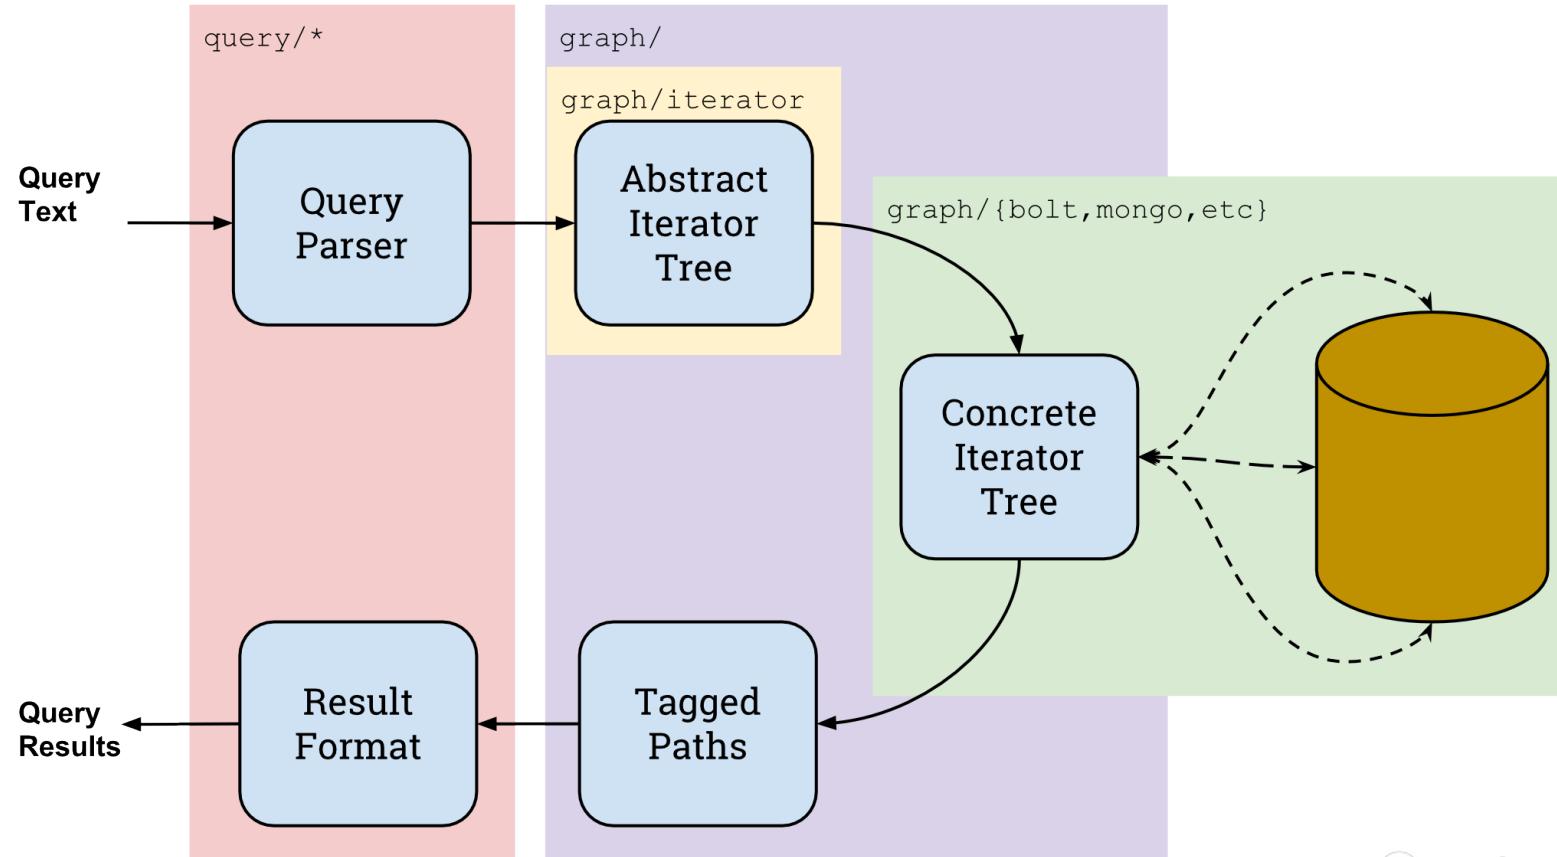
3c.

Quadstores and Concrete Iterators

aka, Backends

graph/{bolt,mongo,leveldb,memstore}

Life of a Graph Query



Role: Store the Data

- Translate quad operations to the particular backend
- Translate opaque database identifiers
- Optimize the query based on backend features

Most translation happens in backend iterators



Primary Interface:

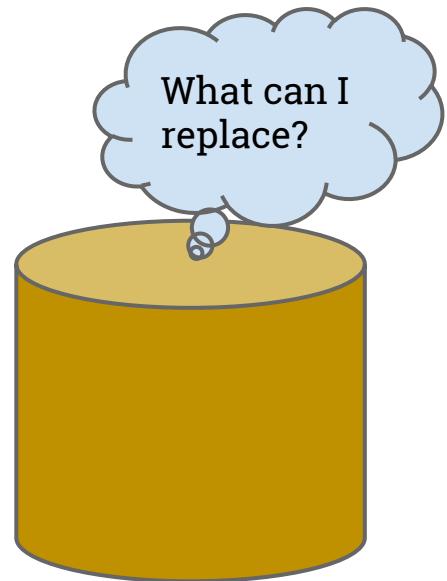
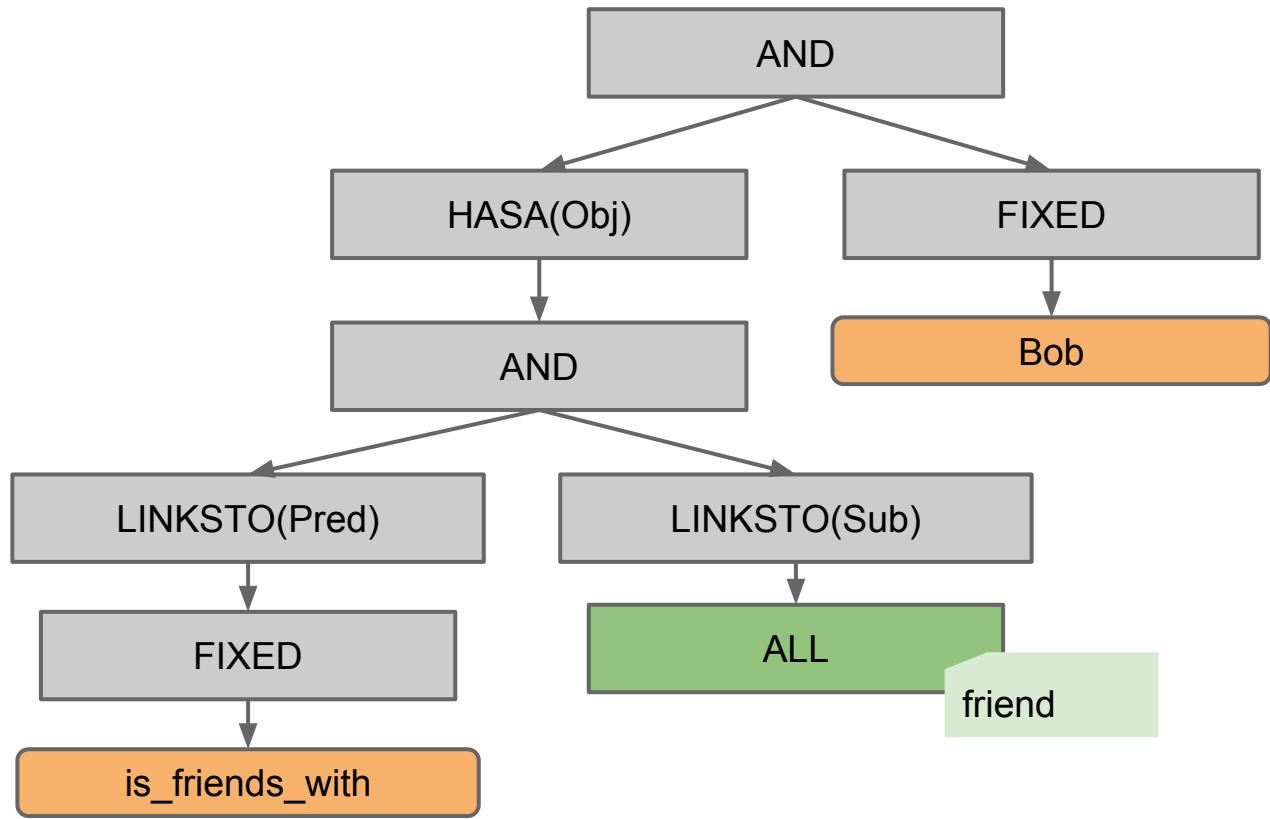
graph/quadstore.go:

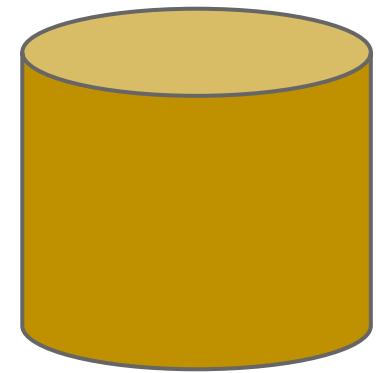
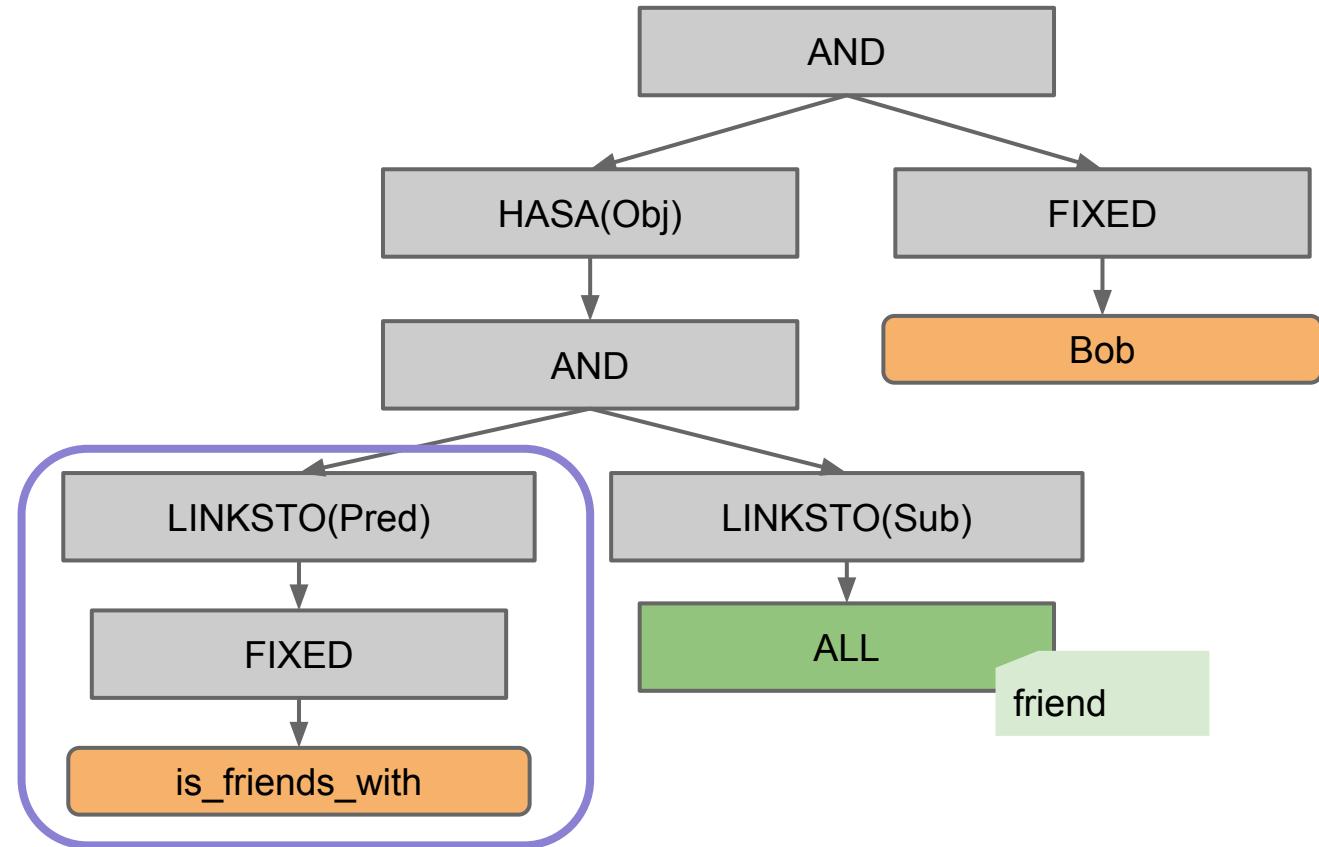
```
// Iterate links that have a given node token in a given field
// For use in HASA and LINKSTO
QuadIterator(quad.Direction, Value) Iterator

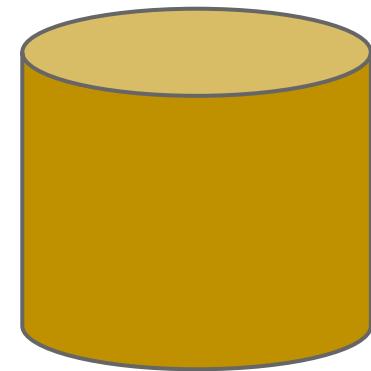
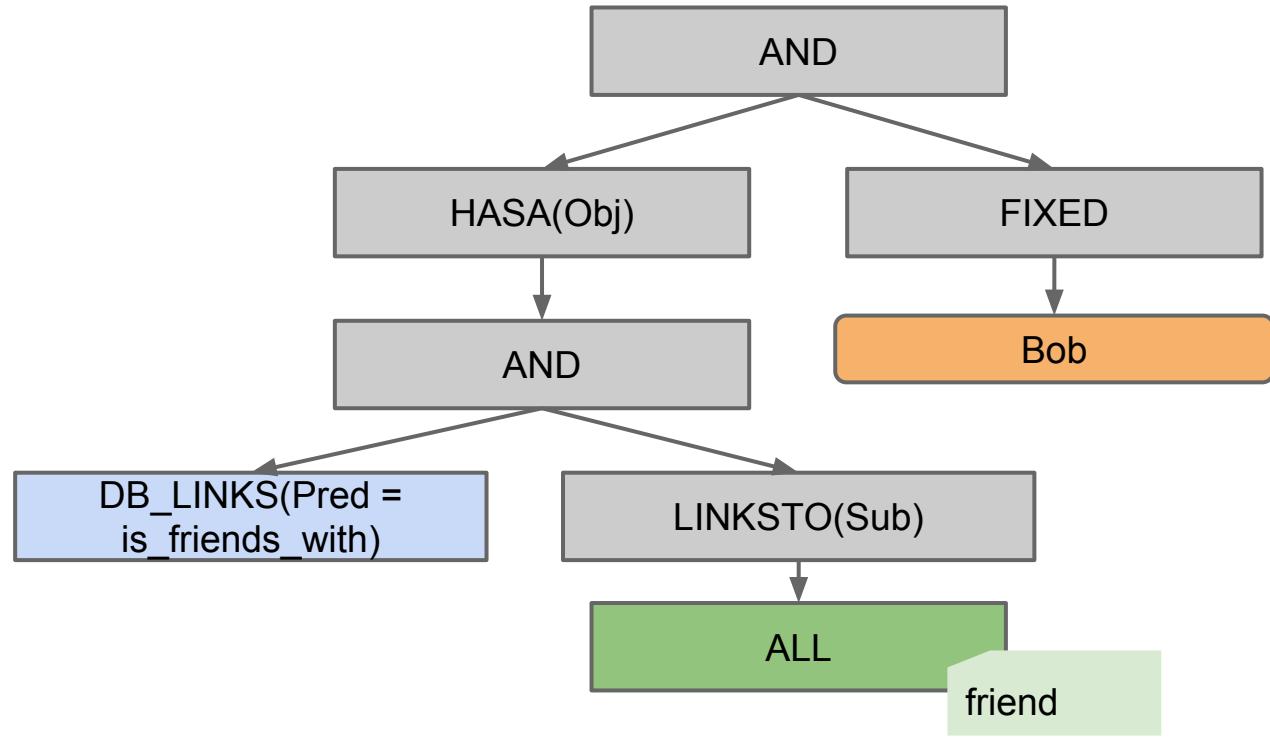
// Returns an ALL iterator for this particular datastore
NodesAllIterator() Iterator

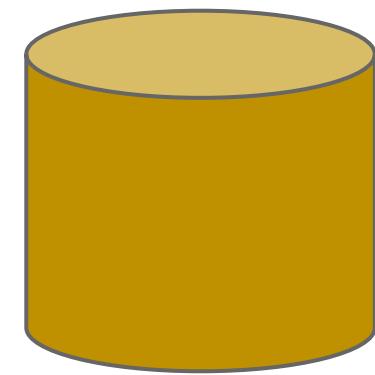
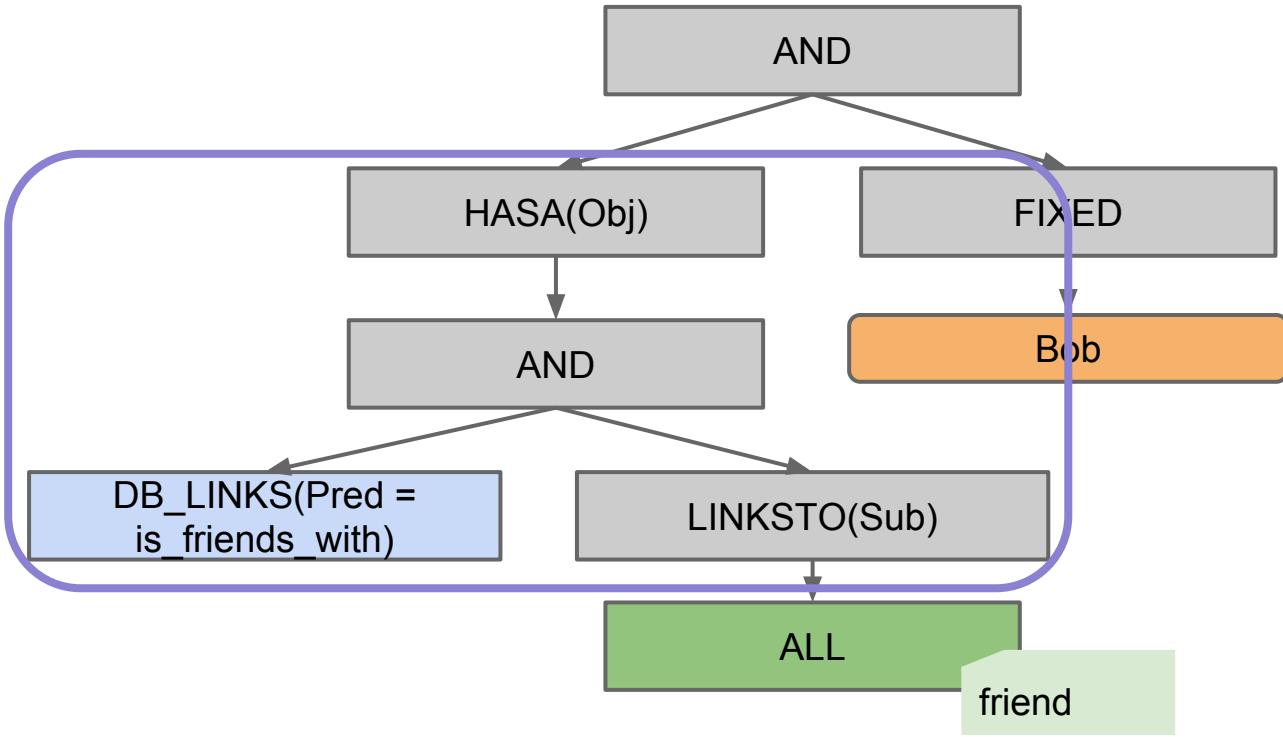
// Make optimization decisions (via replacement)
OptimizeIterator(it Iterator) (Iterator, bool)

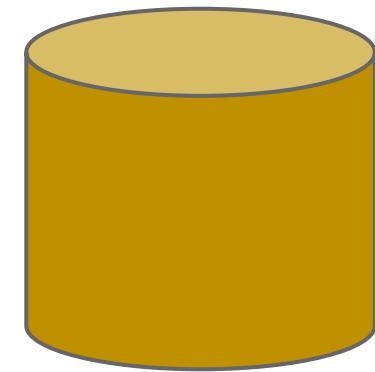
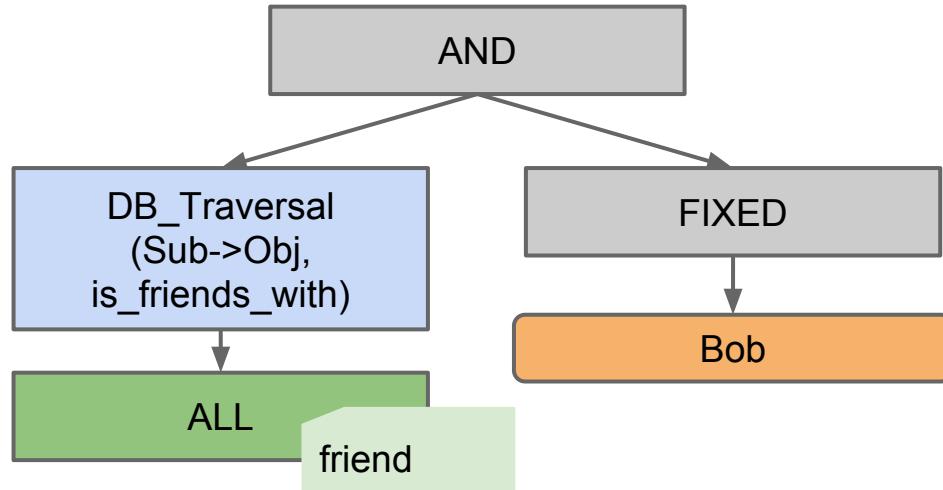
        // Translate the opaque values
ValueOf(string) Value
NameOf(Value) string
Quad(Value) quad.Quad
```







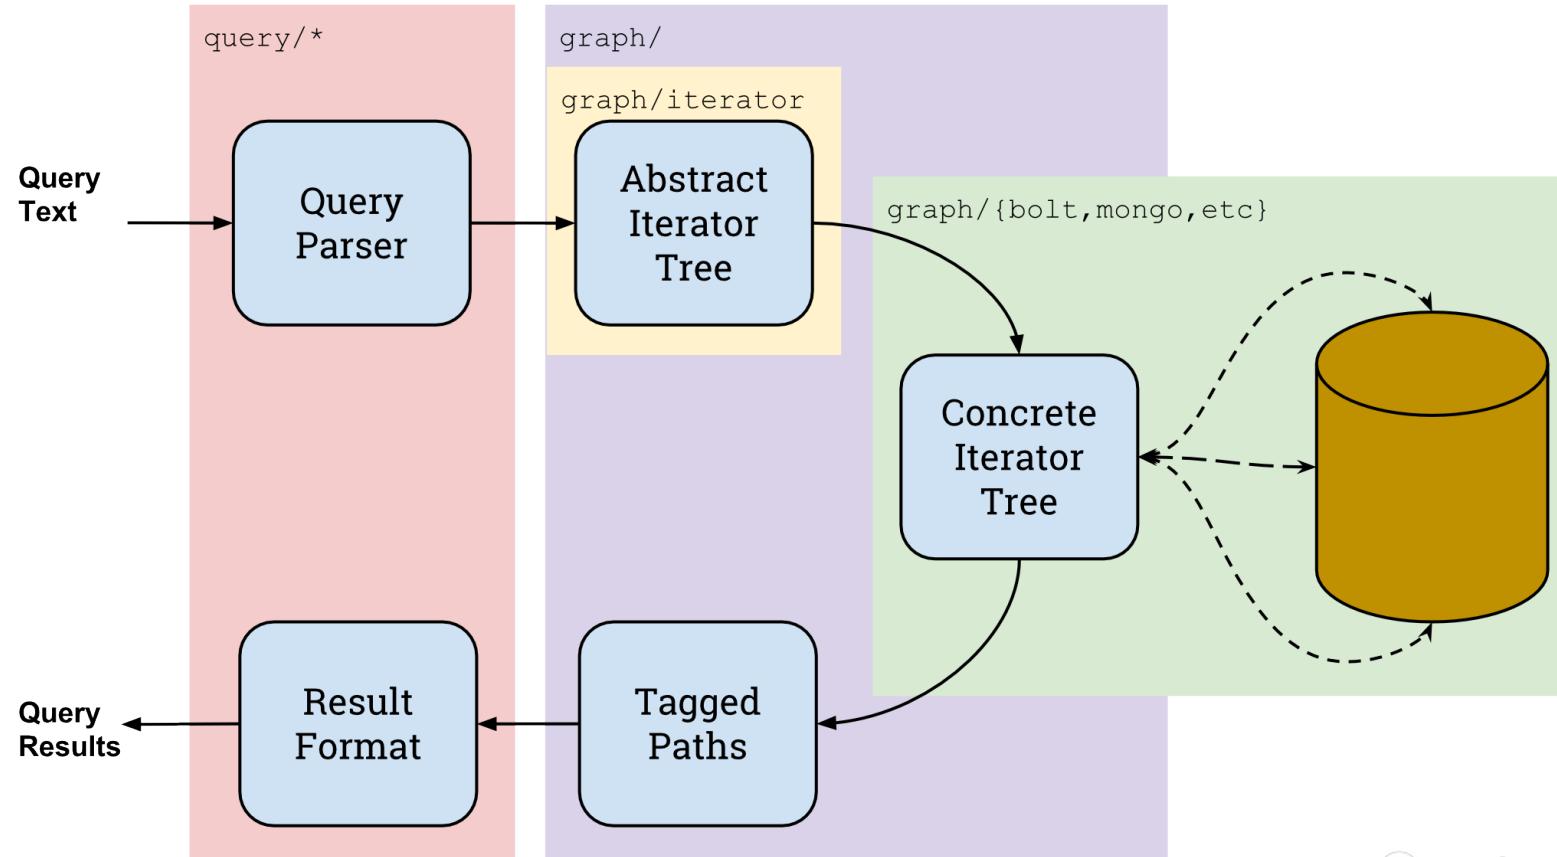


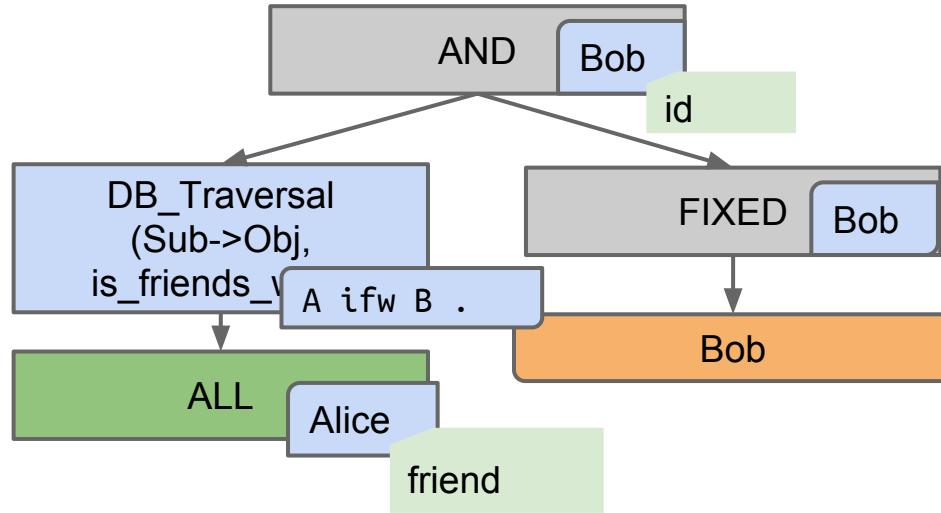


Concrete Iterator Tree

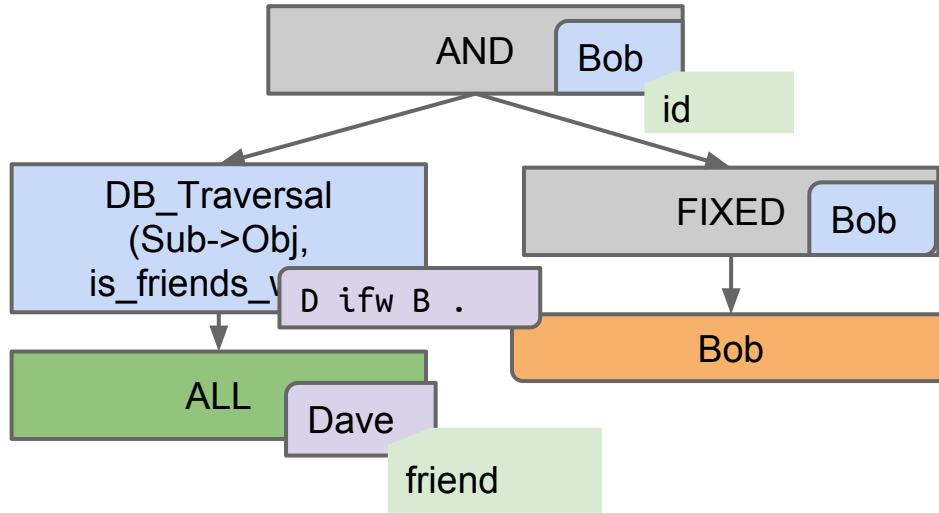


Life of a Graph Query





```
map[string]string{  
    "id": "Bob",  
    "friend": "Alice",  
}
```



```
map[string]string{  
    "id": "Bob",  
    "friend": "Dave",  
}
```

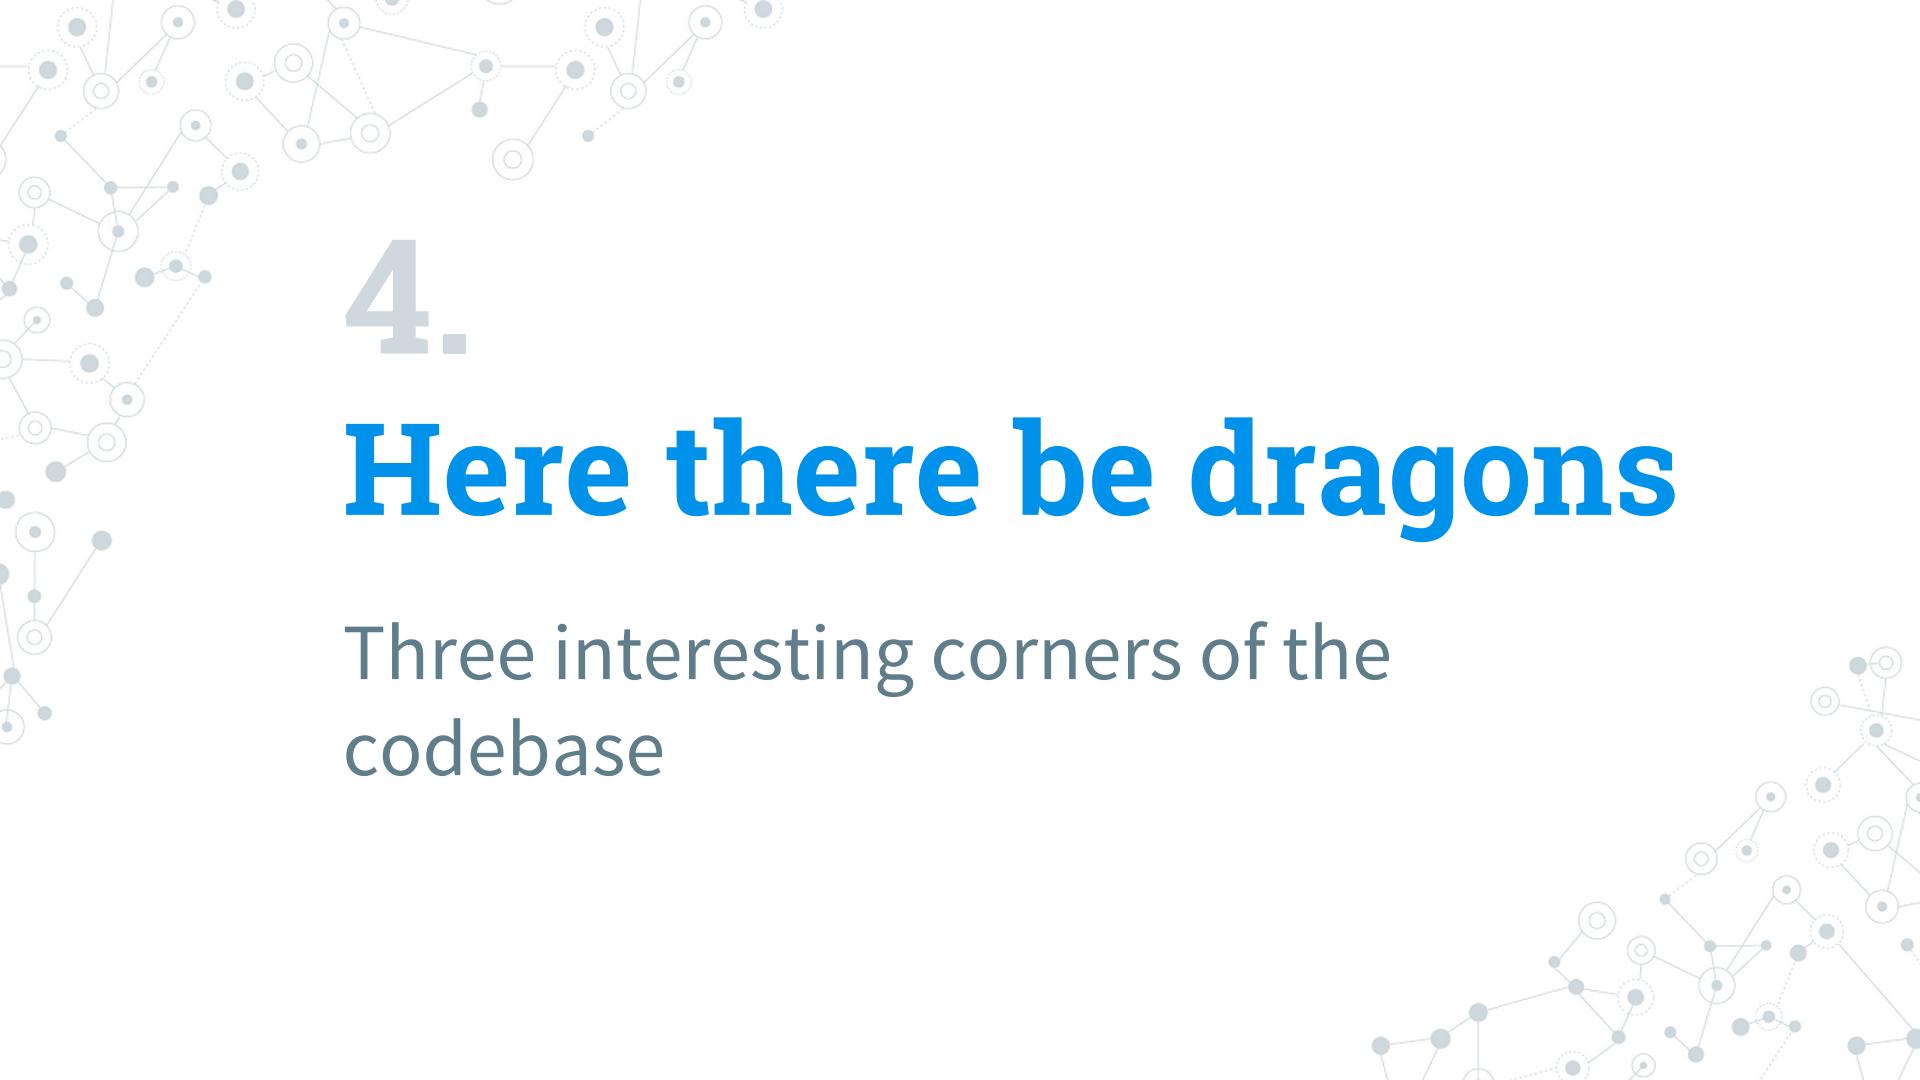
Input:

```
graph.Vertex()  
  .Tag("friend")  
  .Out("is_friends_with")  
  .Is("Bob")  
  .All()
```

Output:

```
{  
  "result": [  
    {  
      "friend": "Alice",  
      "id": "Bob"  
    },  
    {  
      "friend": "Dave",  
      "id": "Bob"  
    }  
  ]  
}
```





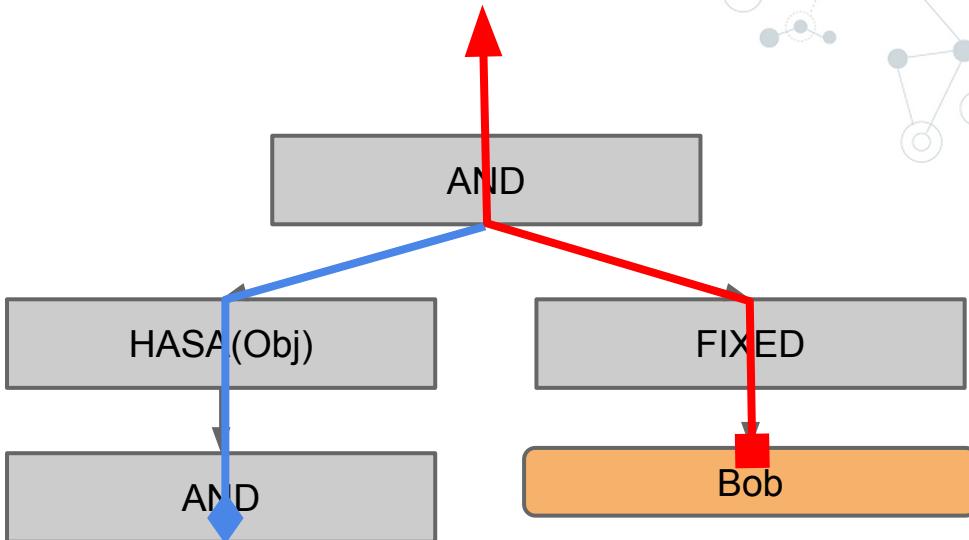
4.

Here there be dragons

Three interesting corners of the codebase

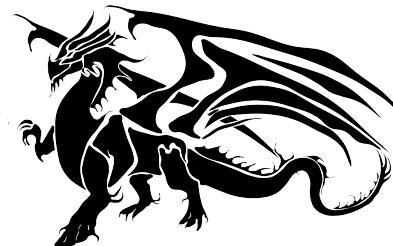
and_iterator_optimize.go

Choosing the best
tree branches to
extract with `Next()`
or check with
`Contains()`



HasA.NextContains()

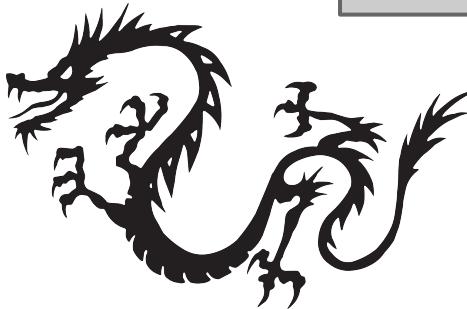
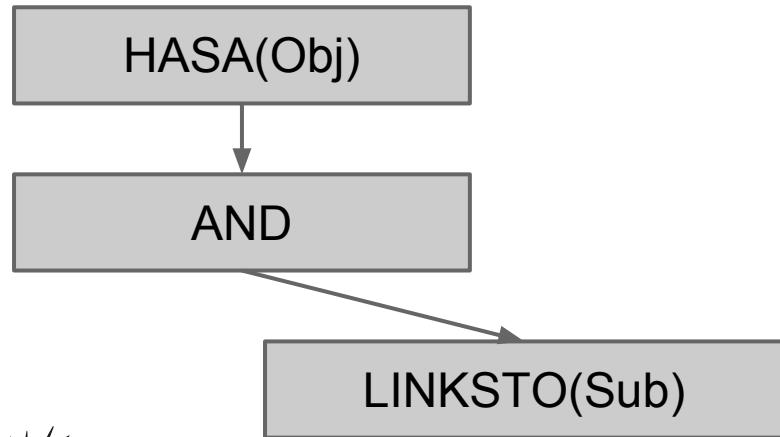
HASA iterators are the potential fan-in point; keeping track of alternate paths is important and subtle



buildIteratorTreeHelper()

.Out()

Defines the tree patterns built for each of the various graph methods exposed in Javascript



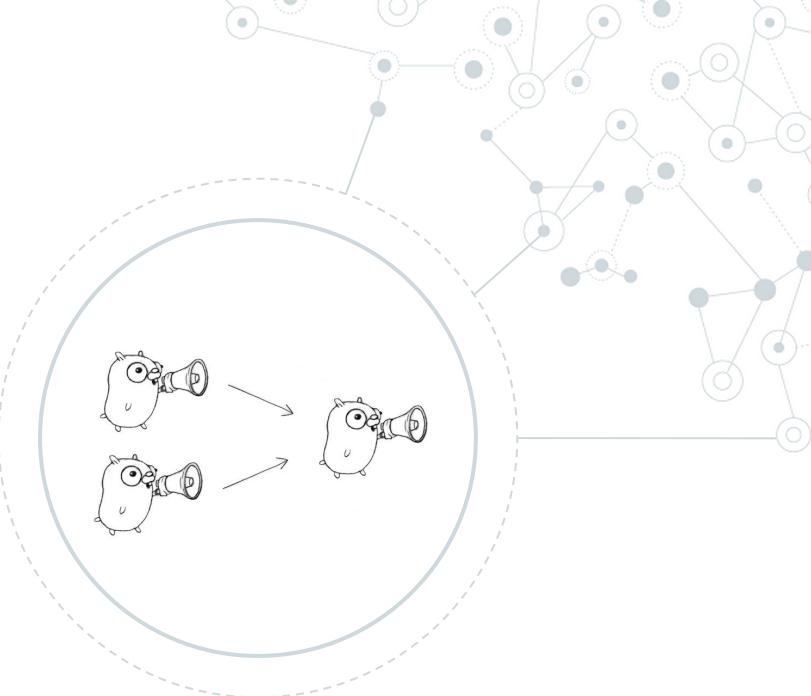
5.

Future Work

Building an even better graph

Parallel Queries

Evaluate two
(reasonably sized)
query branches
simultaneously



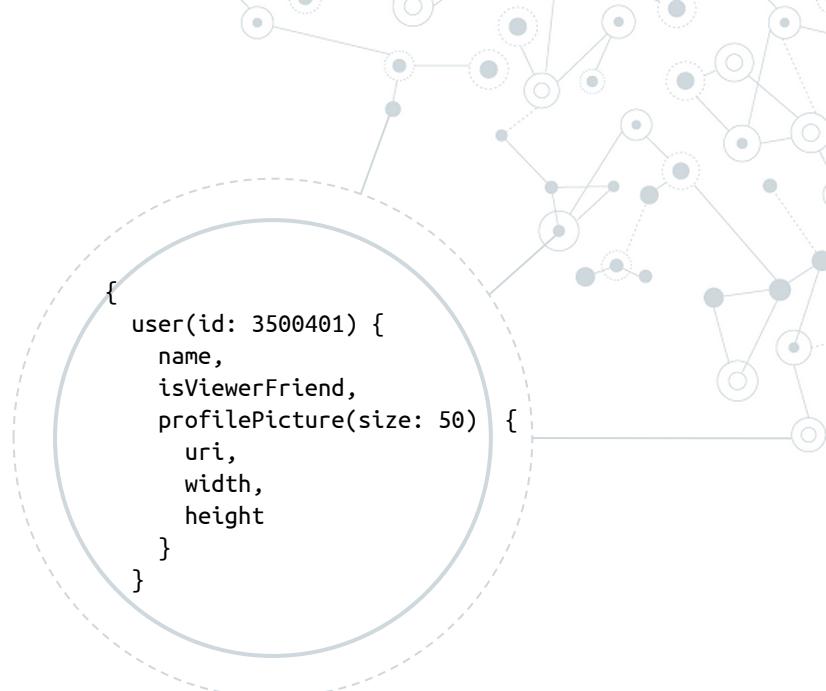
“Perfect” Set Sizes

Precompute/estimate
some cardinalities via
HyperLogLog



New Query Languages

SPARQL, Facebook's
GraphQL, etc.



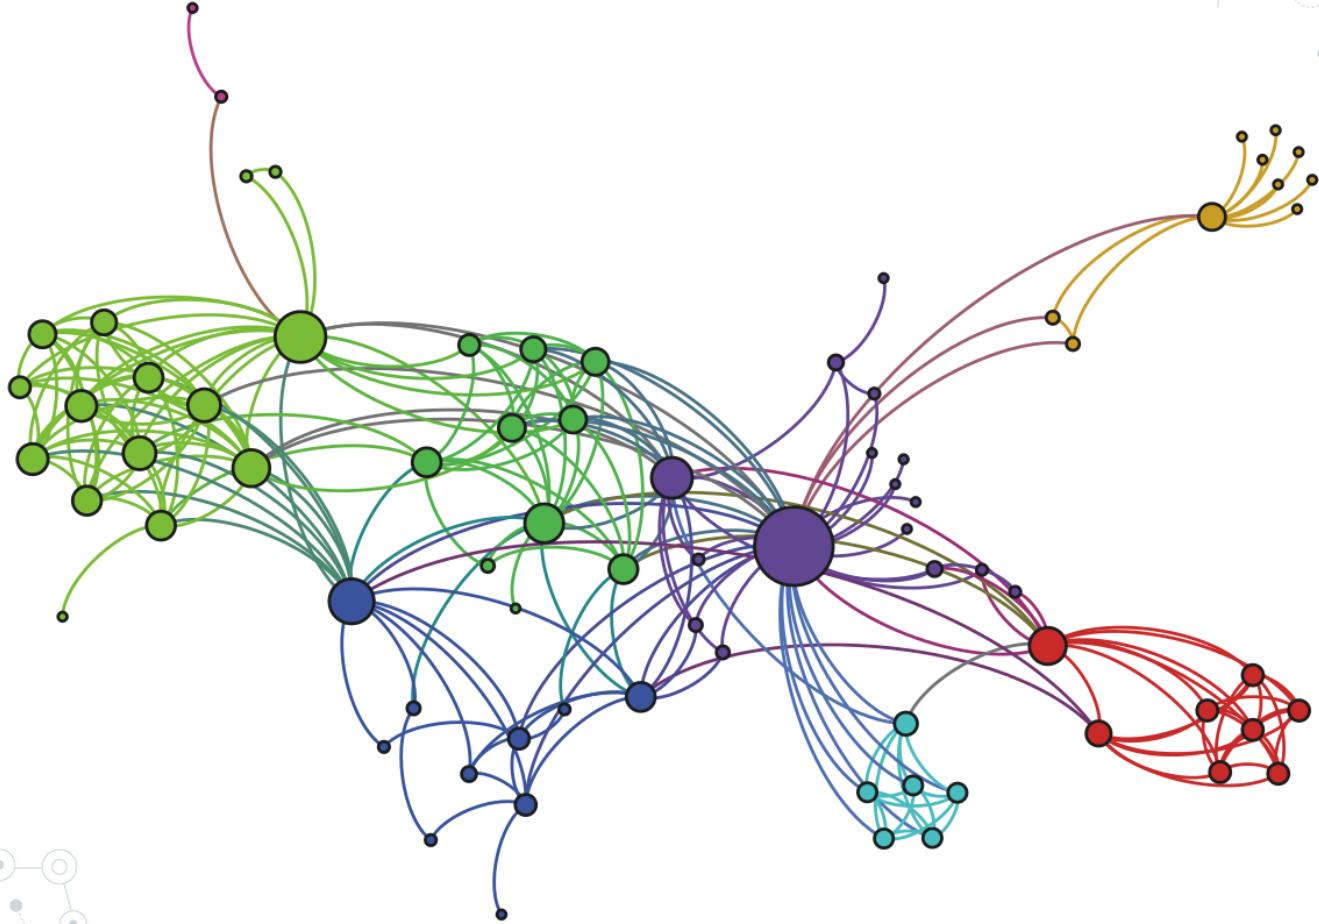
Optional Type Safety

Metadata links declared
on predicates



Many more useful things...

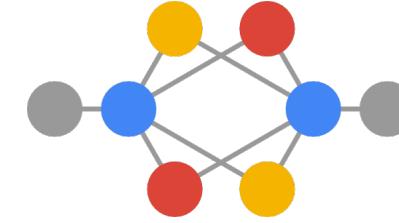
- Cassandra backend (PR for review)
- Full text indexing
- Graph-shaped iterators (generalizing trees)
- Single-master replication
- Longer term: experimental indexes and sharding



"Les Misérables" Social Graph



Questions?



github.com/google/cayley

@barakmich
@cayleygraph

Credits

Special thanks to all the people who made and released these awesome resources for free:

- Presentation template by [SlidesCarnival](#)
 - Alice, Bob and Dave by [Randall Munroe](#)
 - [Kevin Bacon](#) by SAGIndie (via Wikipedia)
 - [Server Icons](#) by RRZEIcons (via Wikipedia)
 - [Connections](#) by Langwitches (via Flickr)
 - [Hydrogen](#) and [Water](#) (via Wikipedia)
 - [square-peg-round-hole-21](#) by Yoel Ben-Avraham (Flickr)
- All dragons CC0