

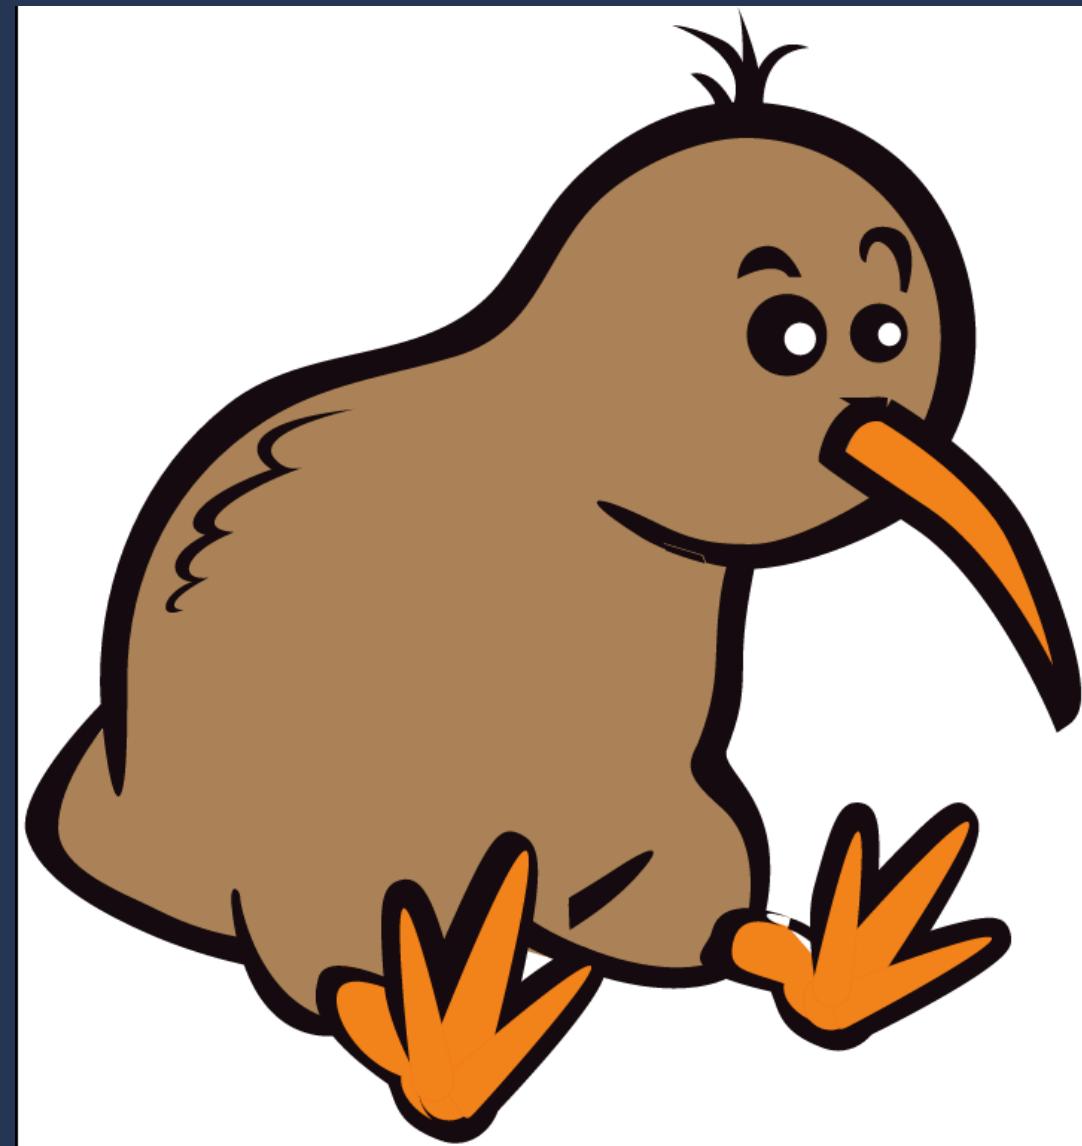
Uptime: Building Resilient Services in Go

Blake Caldwell

<http://blakecaldwell.net>

<https://twitter.com/blakecaldwell>

Fog Creek Software



Git & Hg Source Code Hosting



Kilo

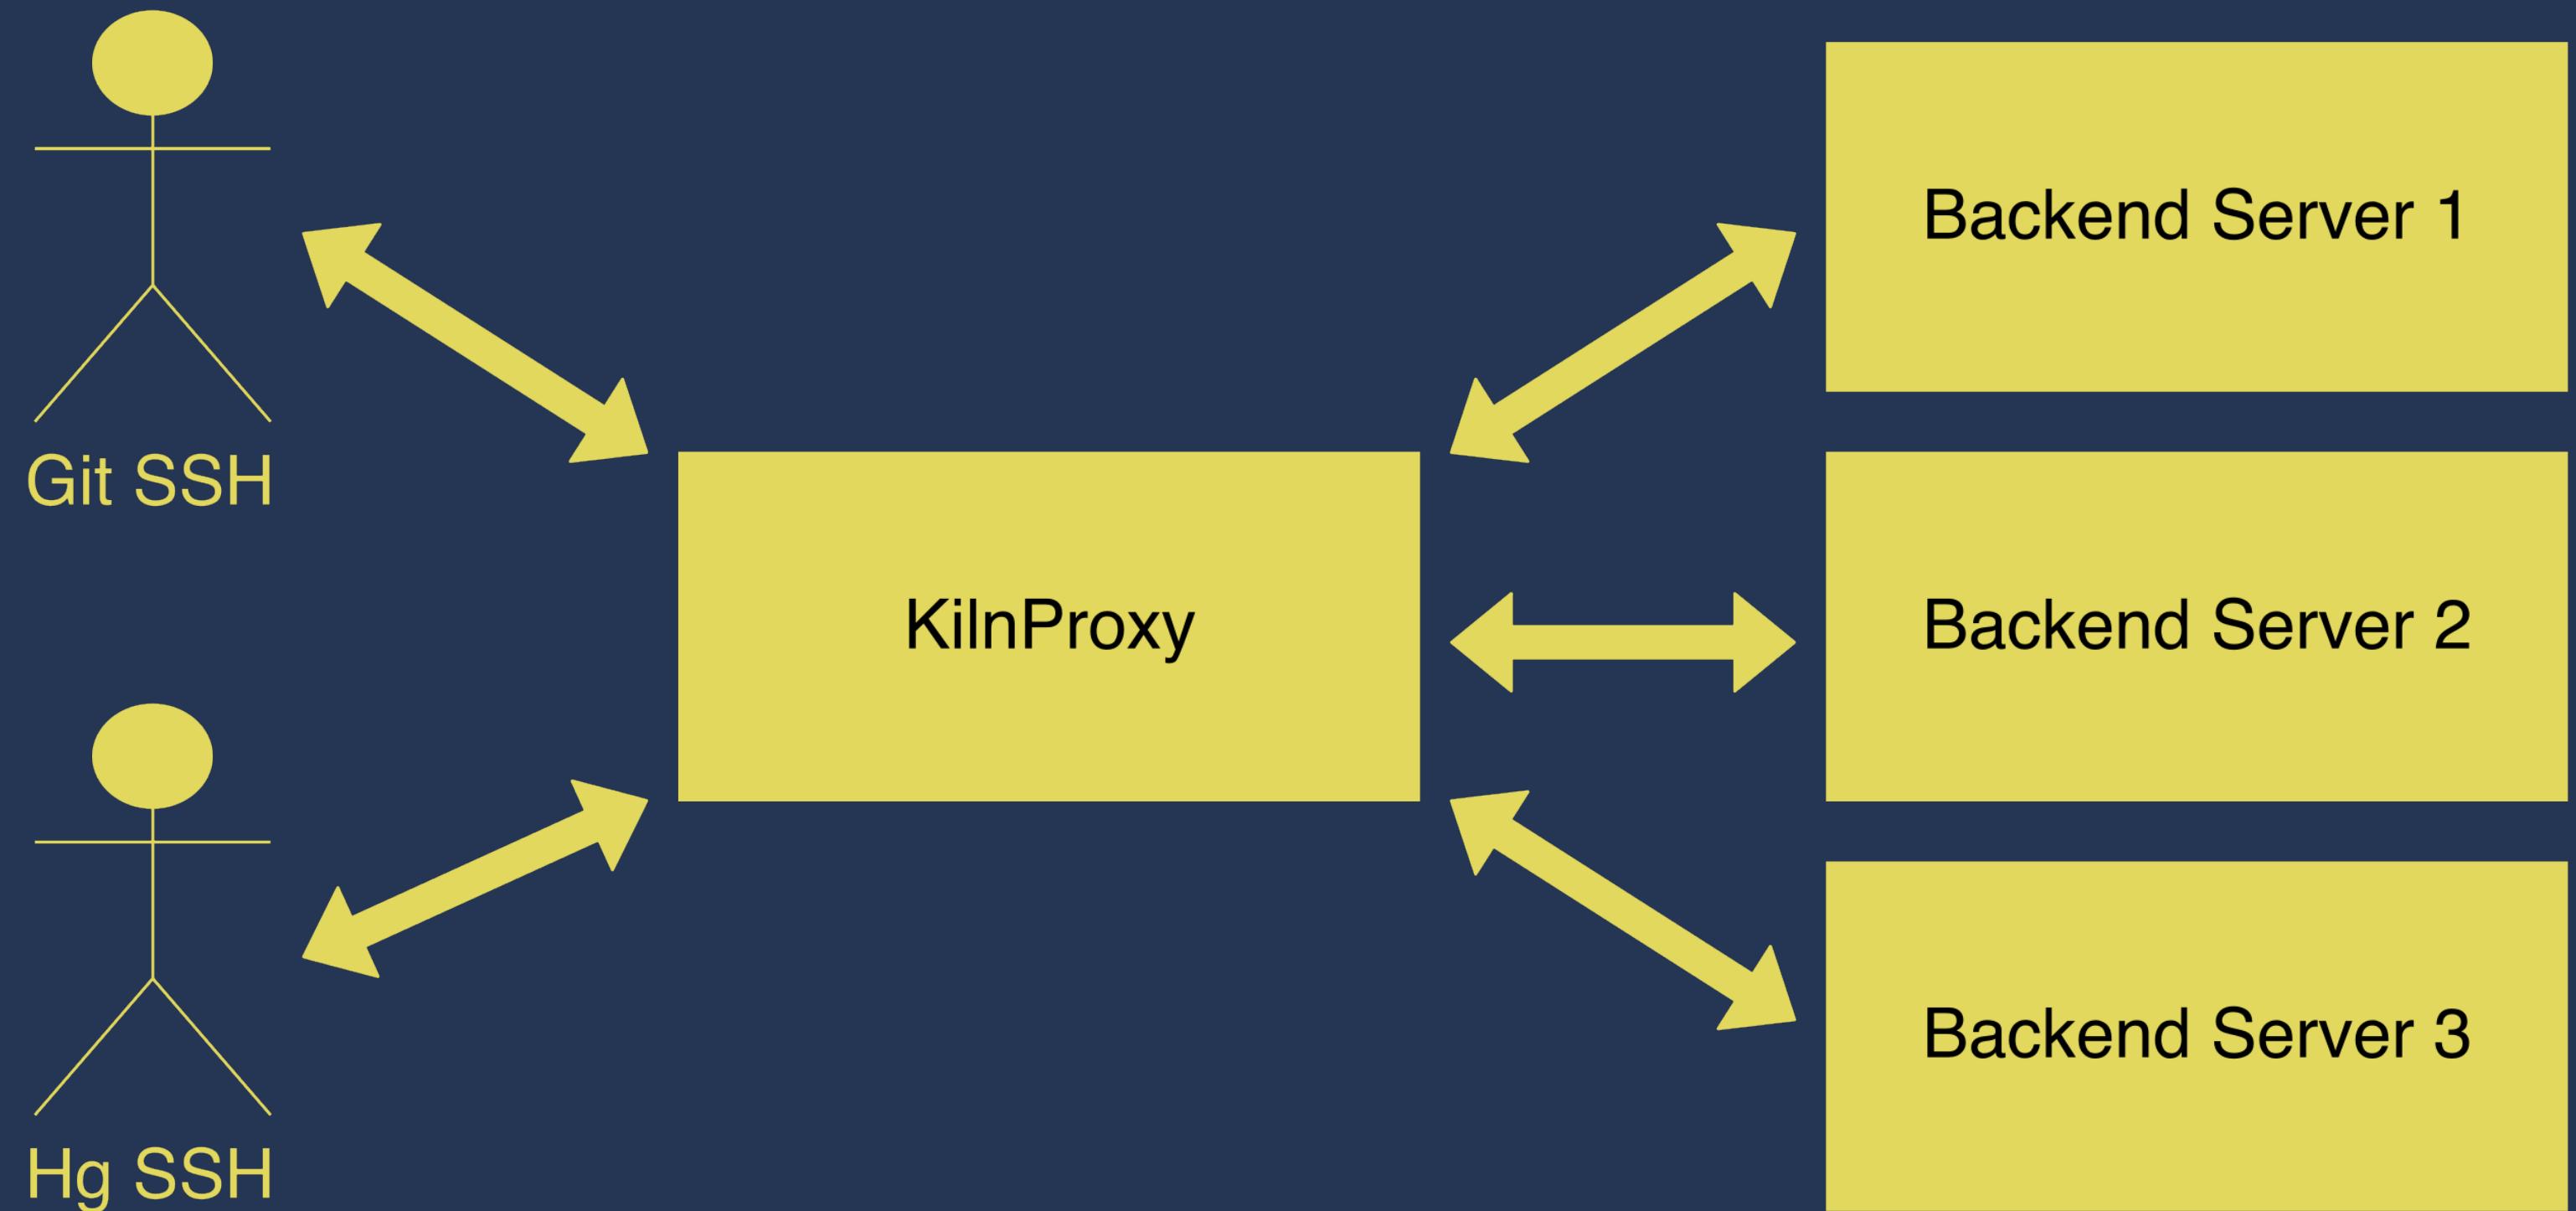
2014 Google I/O: What's Go?



Solution Found!

***Now to find the problem**

Kin's SSH
Reverse Proxy



Why Rewrite?

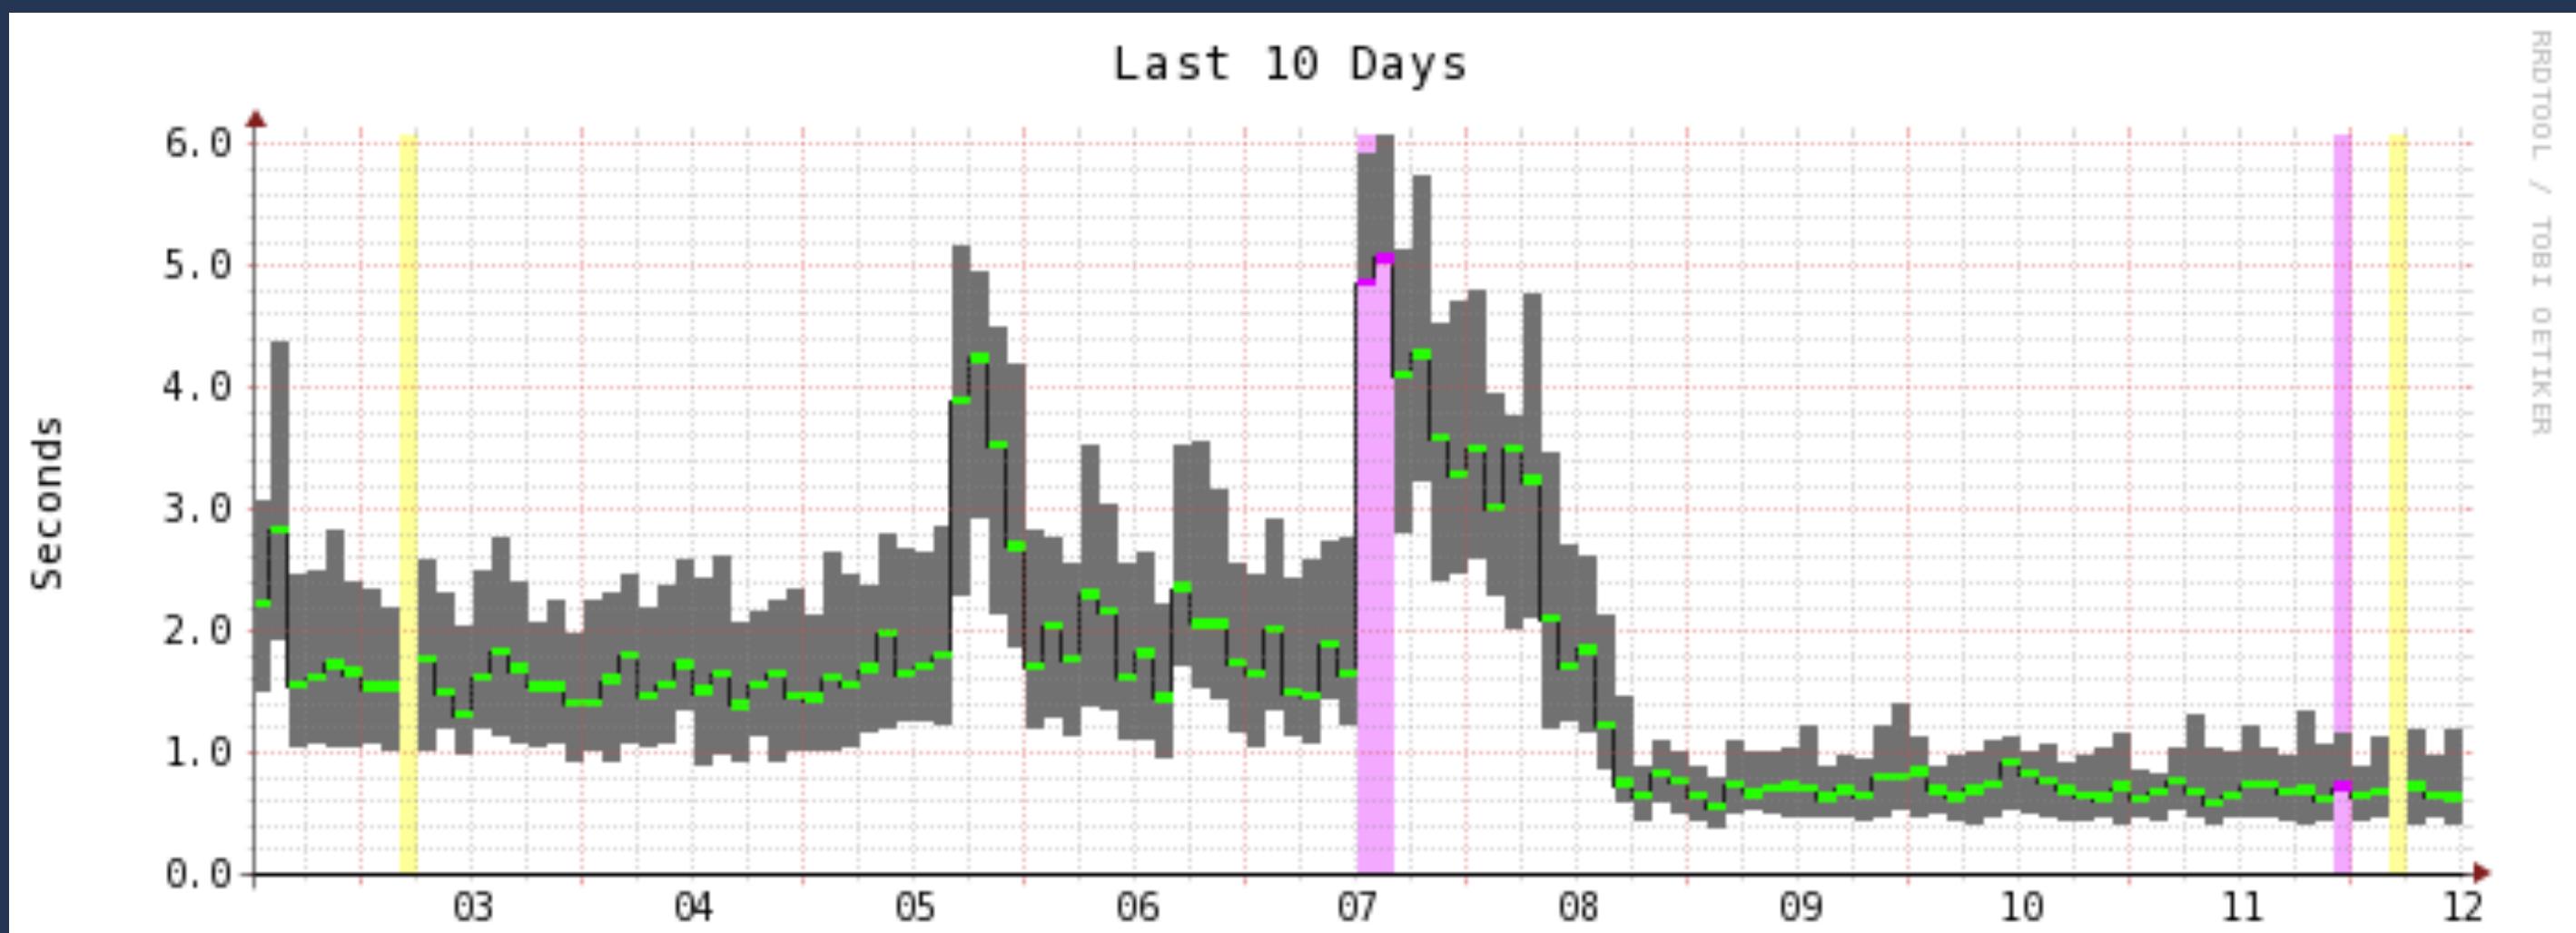
Tons of Concurrency

1. Accepts SSH connection (1 goroutine)
2. Authenticates via public/private key
3. Connects to backend server
4. Proxies STDIN, STDOUT, and STDERR (3 goroutines)

How'd It Go?

KilnProxy: From Python to Go

Clone times halved!



Resiliency: The Process

Careful Coding: Error Handling & Clean-up

Handle All Errors

```
resourceA, err := OpenResourceA()  
if err != nil {  
    return nil, err  
}  
defer resourceA.Close()
```

What About n_1 ?

Back To Our Example:

```
resourceA, err := OpenResourceA()  
if err != nil {  
    return nil, err  
}  
defer resourceA.Close()
```

Be Careful!

```
// can return nil, and that's not an error!
resourceA, err := OpenResourceA()
if err != nil {
    return nil, err
}
defer resourceA.Close() // panic on nil????
```

Not Necessarily.

One Solution:

```
// can return nil, and that's not an error!
resourceA, err := OpenResourceA()
if err != nil {
    return nil, err
}
defer func(){
    if resourceA != nil {
        resourceA.Close()
    }
}()
```

gross

Make deferred Methods nil-Safe

```
// Don't forget: resource might be nil!
func (resource *Resource) Close() {
    if resource != nil {
        // ... clean up
    }
}
```

Much Better!

```
// can return nil, and that's not an error!
resourceA, err := OpenResourceA()
if err != nil {
    return nil, err
}
defer resourceA.Close() // will never panic!
```

Careful Coding: Channels

Channel Axioms ¹

1. A send to a nil channel **blocks forever**
2. A receive from a nil channel **blocks forever**
3. A send to a closed channel **panics**
4. A receive from a closed channel returns the zero value immediately

¹ <http://dave.cheney.net/2014/03/19/channel-axioms>



Careful Coding: Panics!

You Can Recover From Panics

- ... but you shouldn't always do so!
- Only recover if you're sure it's okay
- Panic recovery is for current goroutine
- At very least, log the stack trace

A blurred background image of a Formula 1 race car driving on a track. The car is blue with various sponsor logos, including Michelin, Mobil 1, and Hyatt. The number 18 is visible on the side. The overall motion blur suggests speed.

Careful Coding: Avoid Race Conditions!

Go's Race Detector:

- Reports when variable access is not synchronized
- Crashes with a full stack trace, including the read and write goroutines
- Should be used in unit tests, development, and testing environments

Race Detector Output:

```
=====
```

WARNING: DATA RACE

Read by goroutine 5:

```
  main.func·001()
    race.go:14 +0x169
```

Previous write by goroutine 1:

```
  main.main()
    race.go:15 +0x174
```

Goroutine 5 (running) created at:

```
  time.goFunc()
    src/pkg/time/sleep.go:122 +0x56
  timerproc()
    src/pkg/runtime/ztime_linux_amd64.c:181 +0x189
```

```
=====
```

Enable Race Detection:

```
$ go test -race mypkg          // to test the package  
$ go run -race mysrc.go        // to run the source file  
$ go build -race mycmd         // to build the command  
$ go install -race mypkg // to install the package
```

Careful Coding: Implement Timeouts

Network Timeouts:

- network dial timeout
- network connection inactivity timeout
- total connection timeout

TEST ALL THE THINGS!

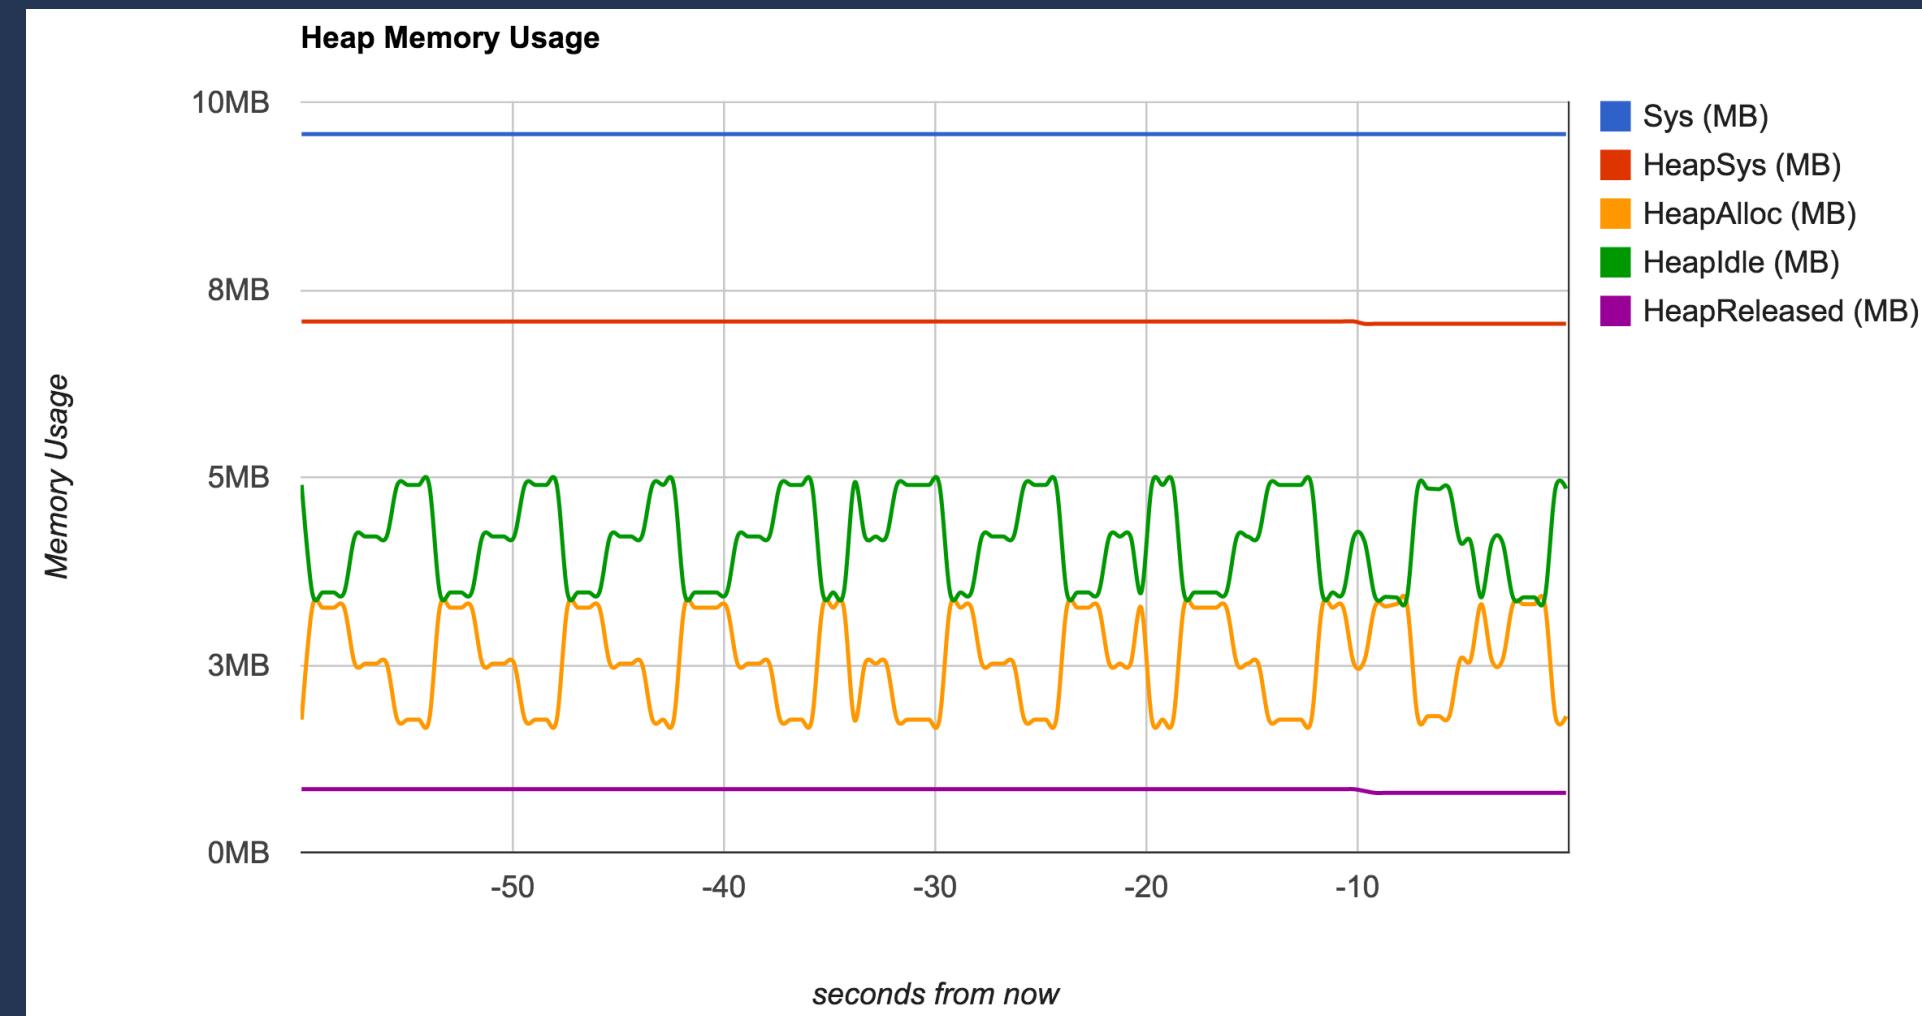


know

Your Service

Know Your Service: How Does It Use Memory?

Profile It!



<http://github.com/fogcreek/profiler>

What to Watch

- How much memory does the service use when idle?
- How much memory per connection?
- Does the system reclaim memory that's no longer used?
- What's the garbage collector doing? `GODEBUG=gctrace=1`
- Where is memory allocated? (PPROF)

Know Your Service: PROF

PPROF Inspects Your Running Process

- Blocking profile
- Goroutine count and full stacktraces
- Heap profile
- Stacktraces that lead to thread creations

Enabling PPROF:

```
import (
    _ "net/http/pprof"
    "net/http"
)

func main() {
    http.ListenAndServe(":6060", nil)
    // ...
}
```

PPROF Main Page

localhost:6060/debug/pprof/

/debug/pprof/

profiles:

- 0[block](#)
- 32[goroutine](#)
- 4[heap](#)
- 8[threadcreate](#)

[full goroutine stack dump](#)

**Don't leak
goroutines!**

Use PPROF To Tell You:

- How many goroutines when nobody is connected?
- How many goroutines per connection?
- Are all goroutines cleaned up after all connections close?

PPROF: Goroutine Page

```
goroutine profile: total 13

1 @ 0x1e5ba6 0x1e594d 0x1e1a35 0xf3edf 0xf404e 0xd24f1 0xd3c5d 0xd462a 0xd2047 0x42e21
#      0x1e5ba6      runtime/pprof.writeRuntimeProfile+0xd6  /Users/blake/go-1.4/src/runtime/pprof/pprof.go:540
#      0x1e594d      runtime/pprof.writeGoroutine+0x9d   /Users/blake/go-1.4/src/runtime/pprof/pprof.go:502
#      0x1e1a35      runtime/pprof.(*Profile).WriteTo+0xd5 /Users/blake/go-1.4/src/runtime/pprof/pprof.go:229
#      0xf3edf       net/http/pprof.handler.ServeHTTP+0x35f /Users/blake/go-1.4/src/net/http/pprof/pprof.go:169
#      0xf404e       net/http/pprof.Index+0x15e        /Users/blake/go-1.4/src/net/http/pprof/pprof.go:181
#      0xd24f1       net/http.HandlerFunc.ServeHTTP+0x41  /Users/blake/go-1.4/src/net/http/server.go:1265
#      0xd3c5d       net/http.(*ServeMux).ServeHTTP+0x17d  /Users/blake/go-1.4/src/net/http/server.go:1541
#      0xd462a       net/http.serverHandler.ServeHTTP+0x19a /Users/blake/go-1.4/src/net/http/server.go:1703
#      0xd2047       net/http.(*conn).serve+0xb57       /Users/blake/go-1.4/src/net/http/server.go:1204

1 @ 0x1a665 0x1a6d8 0x1d30f 0x1d09d 0xf04f9 0x42cb 0x1a3a3 0x42e21
#      0xf04f9 sync.(*WaitGroup).Wait+0x169    /Users/blake/go-1.4/src/sync/waitgroup.go:132
#      0x42cb main.main+0x7b                  /Users/blake/Documents/Development/Personal/src/github.com/wblakecaldwell/sump/server/main.go:112
#      0x1a3a3 runtime.main+0xf3            /Users/blake/go-1.4/src/runtime/proc.go:63

1 @ 0x1a665 0x1a6d8 0x1a4fe 0x42e21
#      0x1a665 runtime.gopark+0x105      /Users/blake/go-1.4/src/runtime/proc.go:131
#      0x1a6d8 runtime.goparkunlock+0x48  /Users/blake/go-1.4/src/runtime/proc.go:136
#      0x1a4fe runtime.forcegchelper+0xce /Users/blake/go-1.4/src/runtime/proc.go:99

1 @ 0x1a665 0x1a6d8 0x13d2c 0x42e21
#      0x1a665 runtime.gopark+0x105      /Users/blake/go-1.4/src/runtime/proc.go:131
#      0x1a6d8 runtime.goparkunlock+0x48  /Users/blake/go-1.4/src/runtime/proc.go:136
#      0x13d2c runtime.bgsweep+0xbc     /Users/blake/go-1.4/src/runtime/mgc0.go:98

1 @ 0x1a665 0x1a6d8 0x1332a 0x42e21
#      0x1a665 runtime.gopark+0x105      /Users/blake/go-1.4/src/runtime/proc.go:131
#      0x1a6d8 runtime.goparkunlock+0x48  /Users/blake/go-1.4/src/runtime/proc.go:136
#      0x1332a runtime.runfinq+0xba    /Users/blake/go-1.4/src/runtime/malloc.go:727
```

PROF:

From the Command Line

What Are Your GoRoutines Doing?

```
$ go tool pprof ./server http://localhost:6060/debug/pprof/goroutine
```

```
(pprof) top5
```

```
11 of 11 total ( 100%)
Showing top 5 nodes out of 49 (cum >= 1)
      flat  flat%  sum%          cum  cum%
          9  81.82% 81.82%          9  81.82%  runtime.gopark
          1  9.09% 90.91%          1  9.09%  runtime.notetsleepg
          1  9.09% 100%           1  9.09%  runtime/pprof.writeRuntimeProfile
          0    0% 100%           1  9.09%  bufio.(*Reader).ReadLine
          0    0% 100%           1  9.09%  bufio.(*Reader).ReadSlice
```

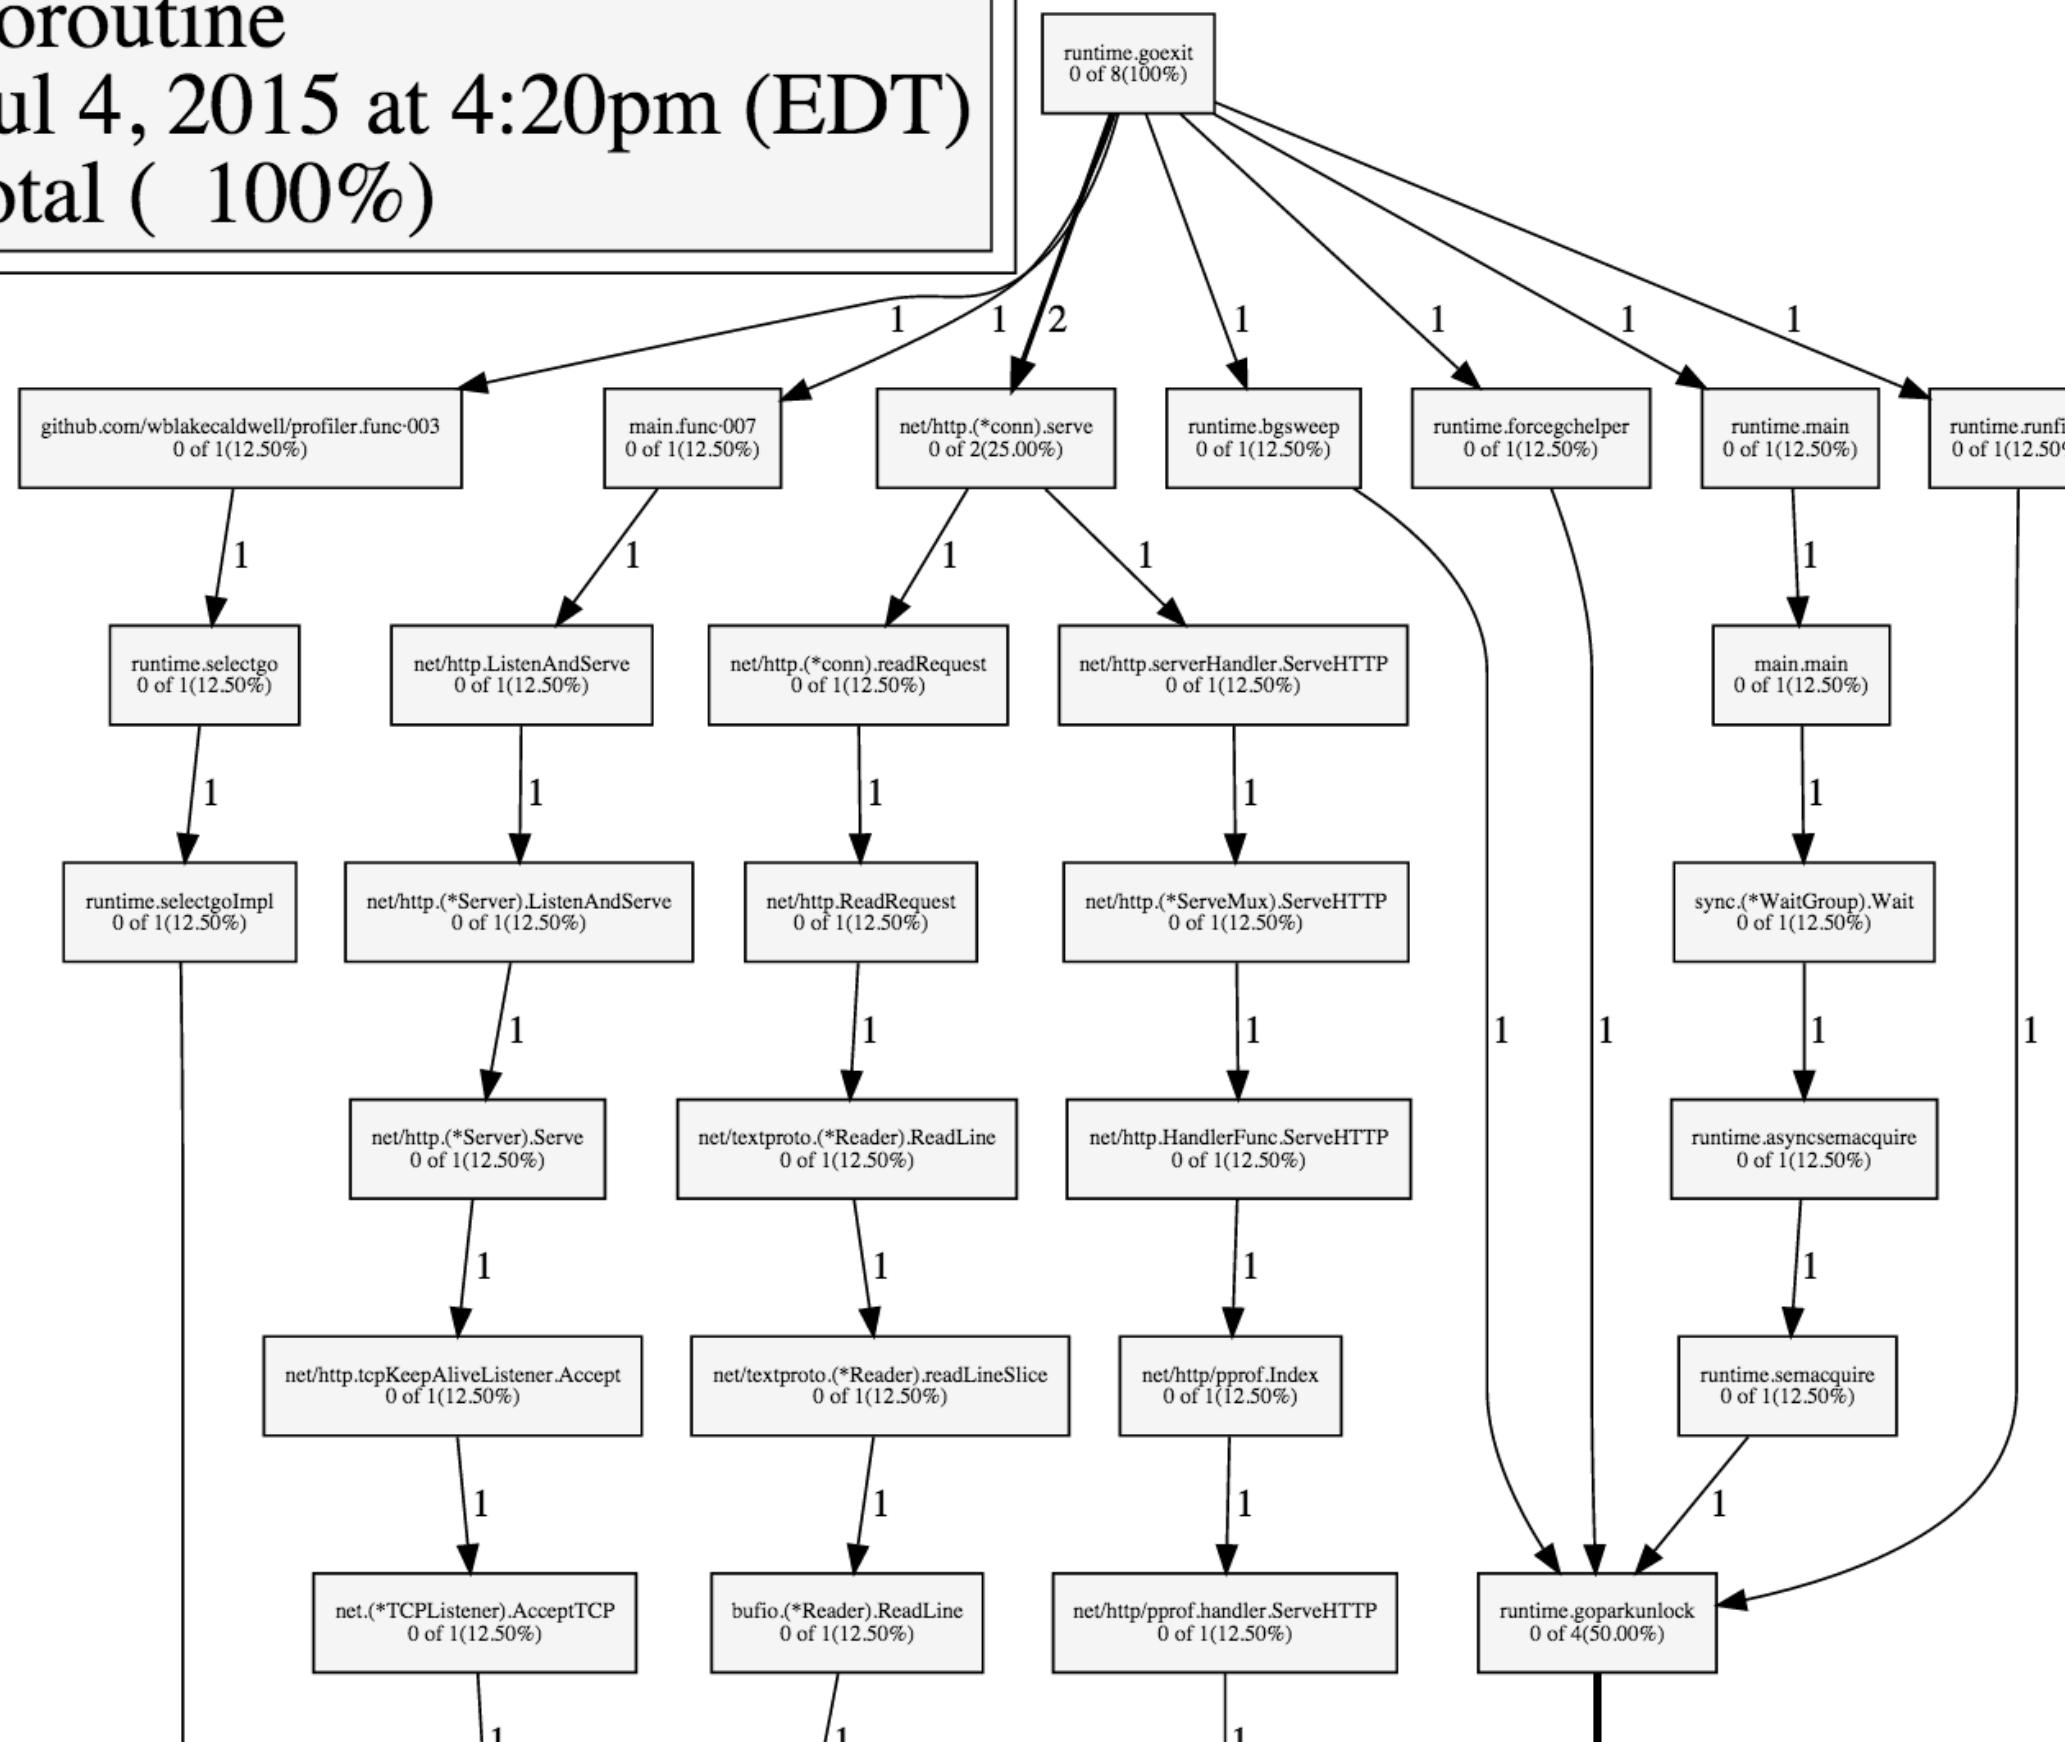
```
(pprof) web
```

File: server

Type: goroutine

Time: Jul 4, 2015 at 4:20pm (EDT)

8 of 8 total (100%)



Who's Allocating Heap Memory?

```
$ go tool pprof ./server http://localhost:6060/debug/pprof/heap
```

```
(pprof) top5
```

```
2362.41kB of 2362.41kB total ( 100%)
```

```
Dropped 28 nodes (cum <= 11.81kB)
```

	flat	flat%	sum%		cum	cum%	
1850.27kB	78.32%	78.32%	78.32%	github.com/wblakecaldwell/profiler.func·002	1850.27kB	78.32%	github.com/wblakecaldwell/profiler.func·002
512.14kB	21.68%	100%	100%	mcommoninit	512.14kB	21.68%	mcommoninit
0	0%	100%	100%	runtime.goexit	1850.27kB	78.32%	runtime.goexit
0	0%	100%	100%	runtime.rt0_go	512.14kB	21.68%	runtime.rt0_go
0	0%	100%	100%	runtime.schedinit	512.14kB	21.68%	runtime.schedinit

```
(pprof) web
```

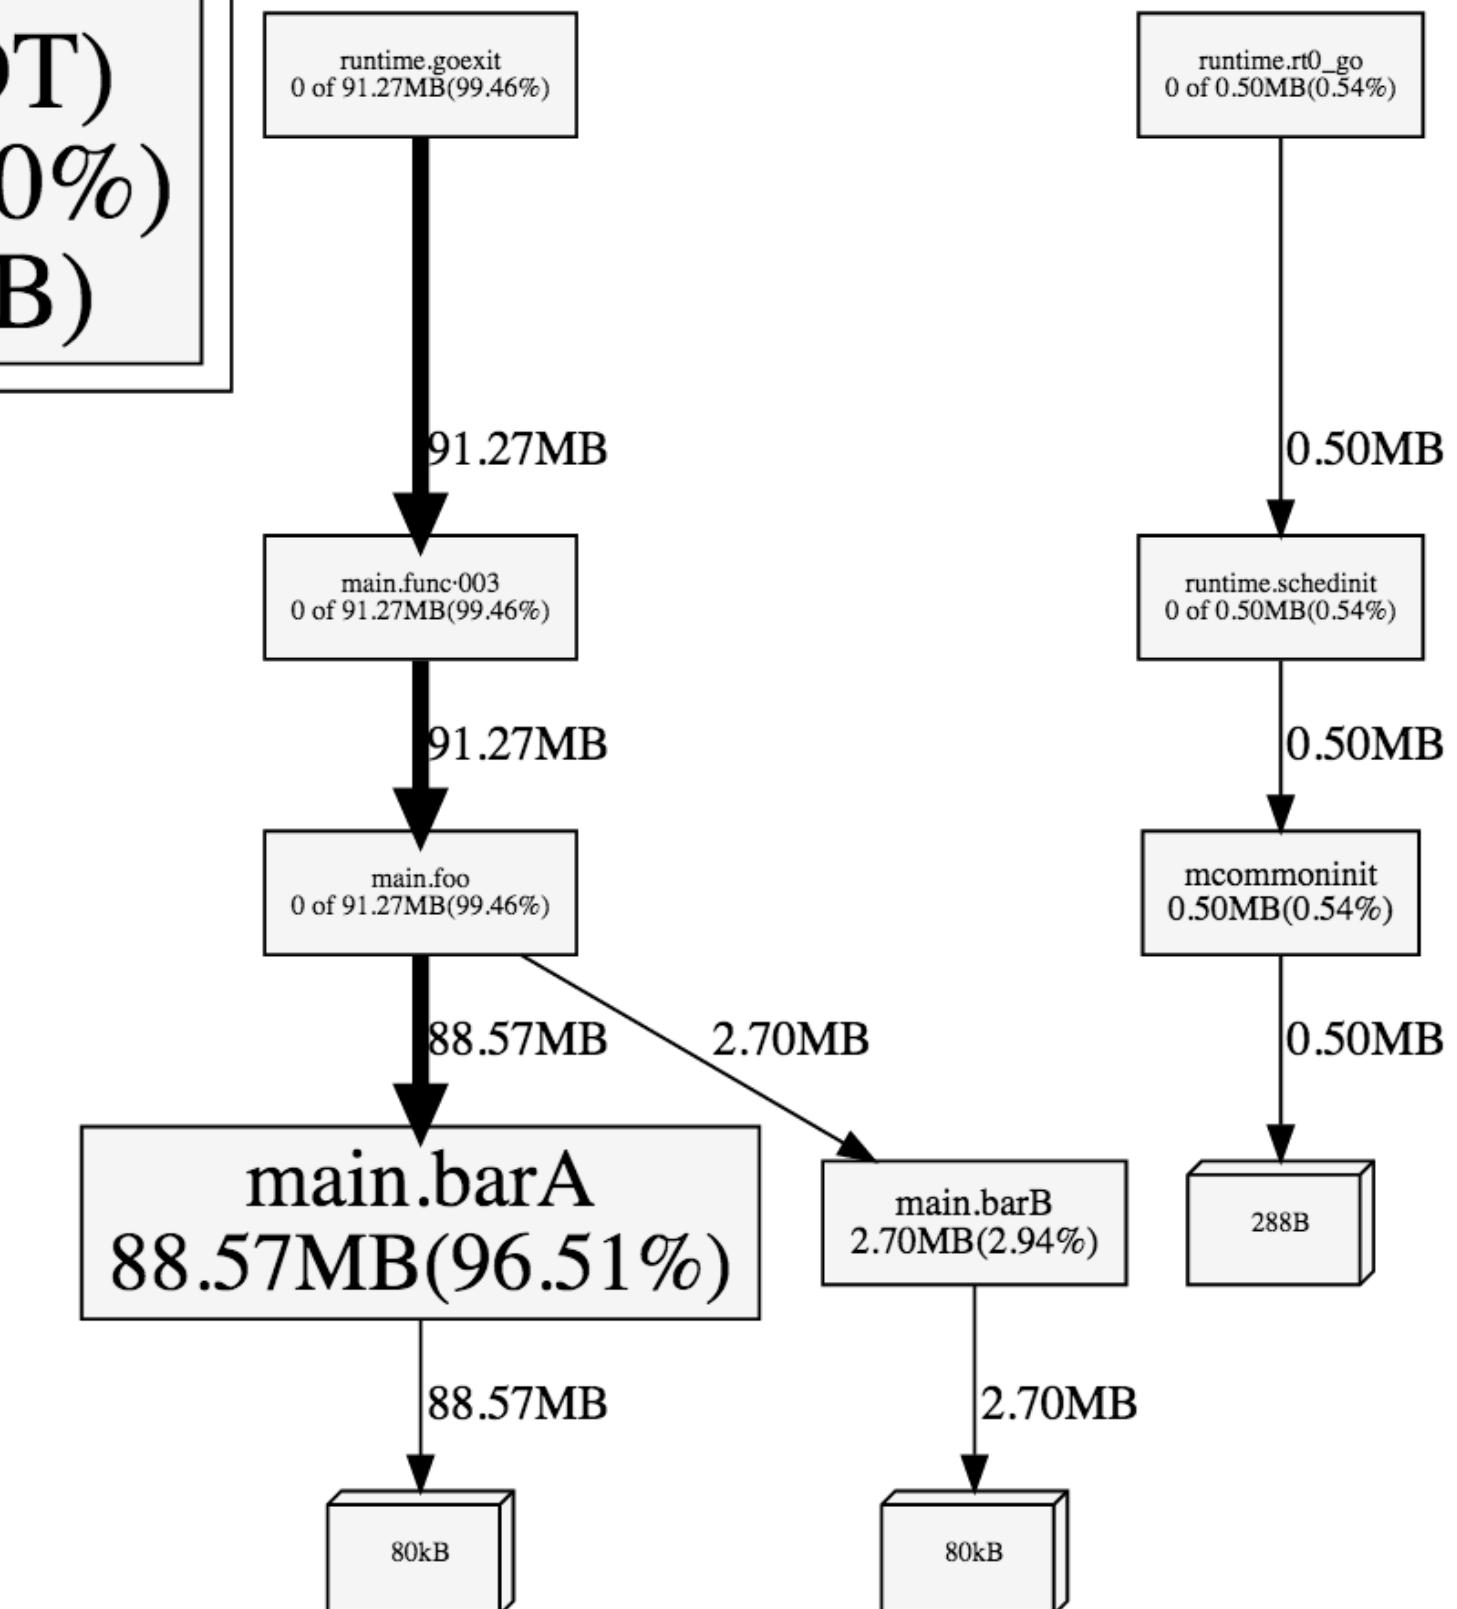
File: server

Type: inuse_space

Time: Jul 4, 2015 at 4:29pm (EDT)

91.77MB of 91.77MB total (100%)

Dropped 1 node (cum <= 0.46MB)



**Know Your Service:
watch it Run**

/info Endpoint

```
{  
    Version: "1.0.275-b244a2b9b8-20150202.163449",  
    StartTimeEpochSecs: 1430515329,  
    CurrentTimeEpochSecs: 143117131,  
    Uptime: "167h10m2s"  
}
```

Managing Service Version

Version: "1.0.275-b244a2b9b8-20150202.163449"

Which is:

<major>. <minor>. <commit#>-<Git SHA>-<date>. <time>

Managing Service Version

Version is stored in a global variable, set by your build script

In code:

```
var ServiceVersion string
```

Build script:

```
$ go build -ldflags \  
  "-X main.ServiceVersion \  
  1.0.275-b244a2b9b8-20150202.163449" \  
  kilnproxy
```

Keep Good Logs!

- Create a semi-unique string per request
- Use this request string as a prefix in all log entries
- Always log at least the start and end of a request

**Who's currently
connected?**

/connections Endpoint

```
{  
  "CurrentUserCount":1,  
  "CurrentlyAuthenticatedUsers":  
  [  
    {"Account":"aviato",  
     "Name":"Erlich Bachman",  
     "PublicKeyName":"Build Server",  
     "SessionKey":"106abc0c",  
     "SessionDuration":"25m4s"  
    }  
  ]  
}
```

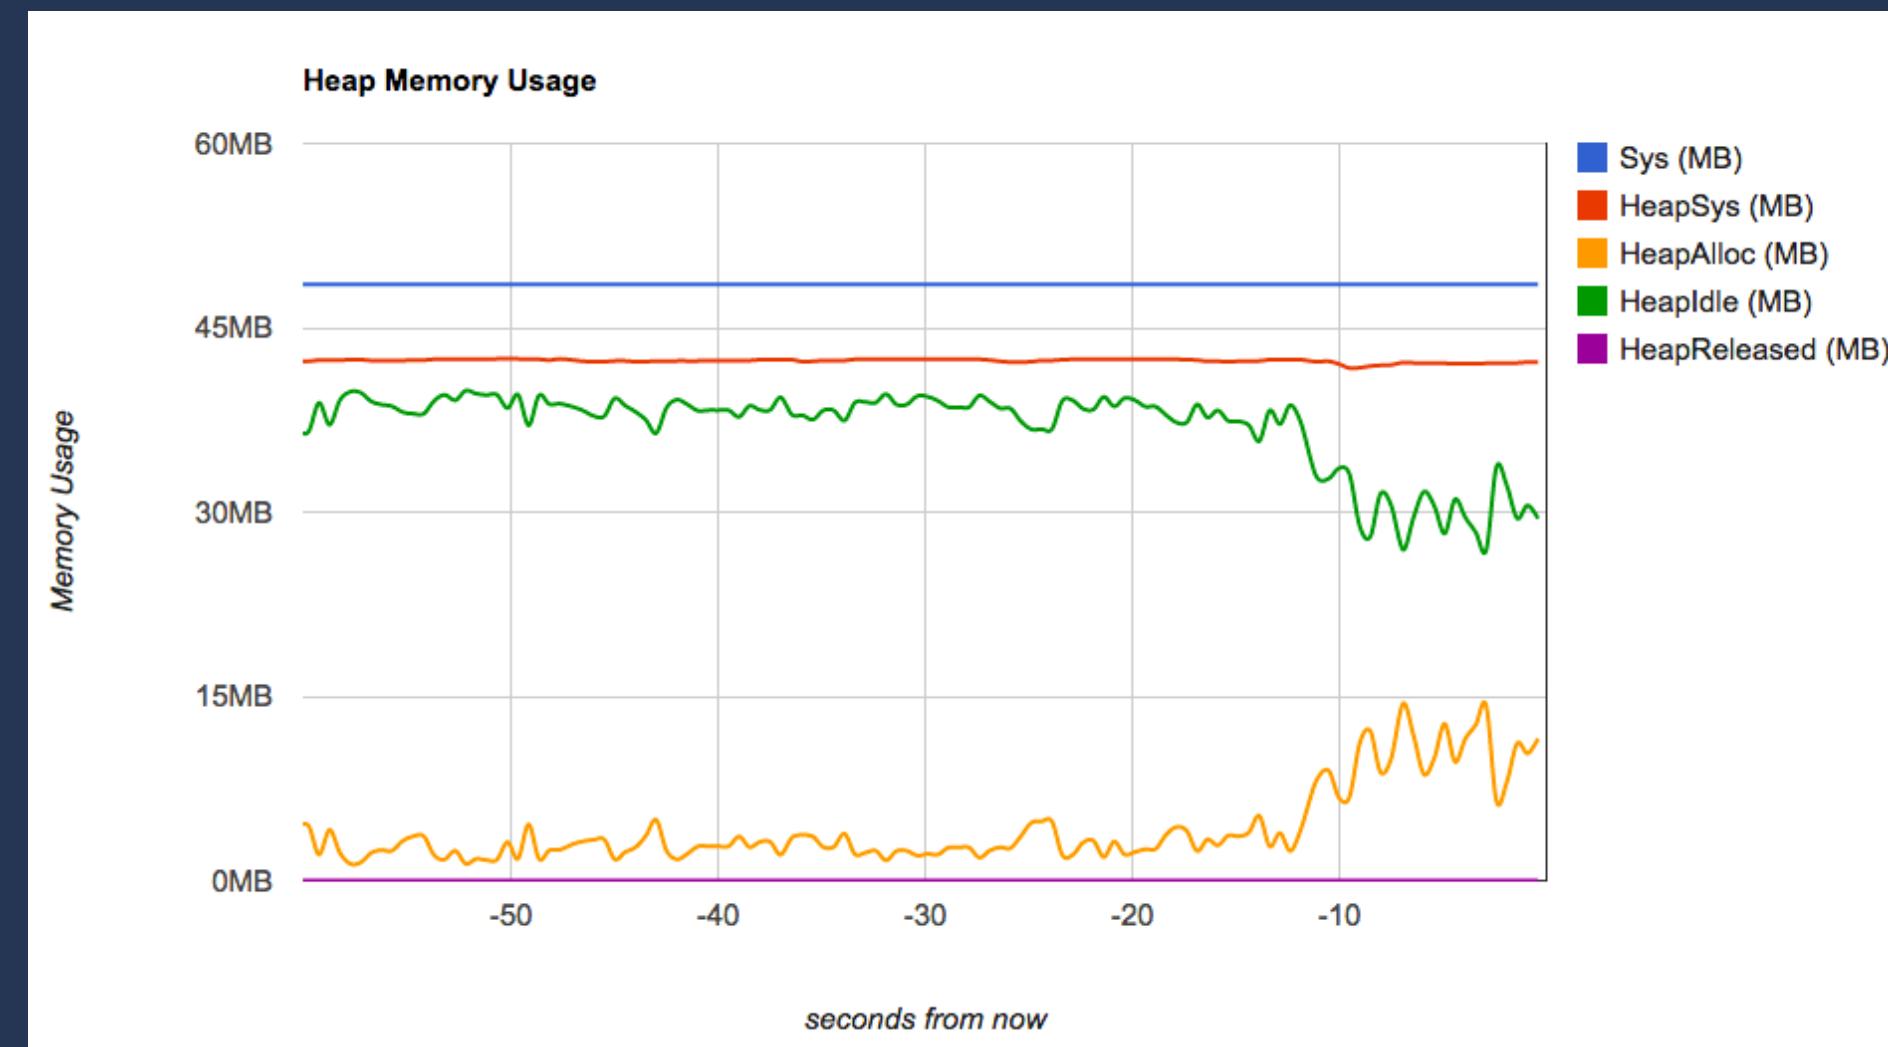
rain

and die

**Game
Day.**

**KilnProxy Using 40MB
more memory than normal!!!**

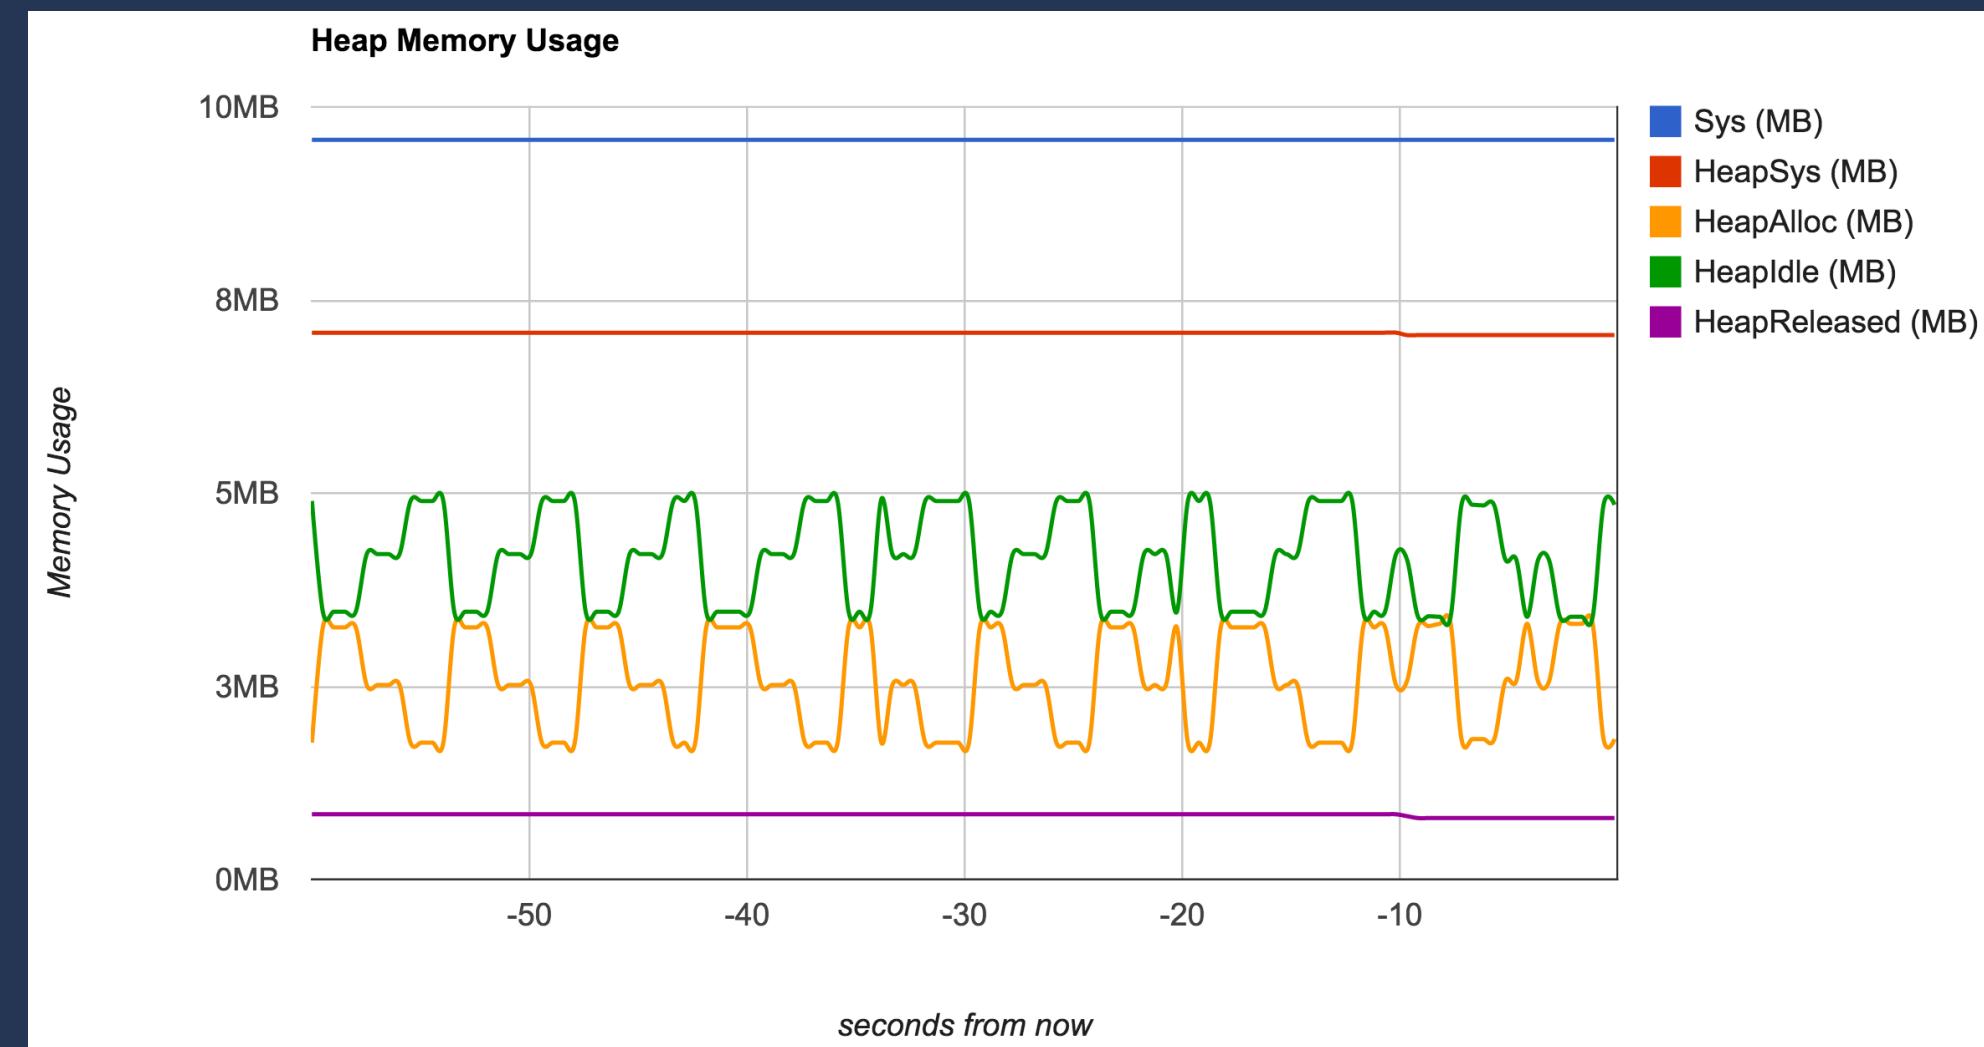
Profiler Tells Me: This memory is still in use



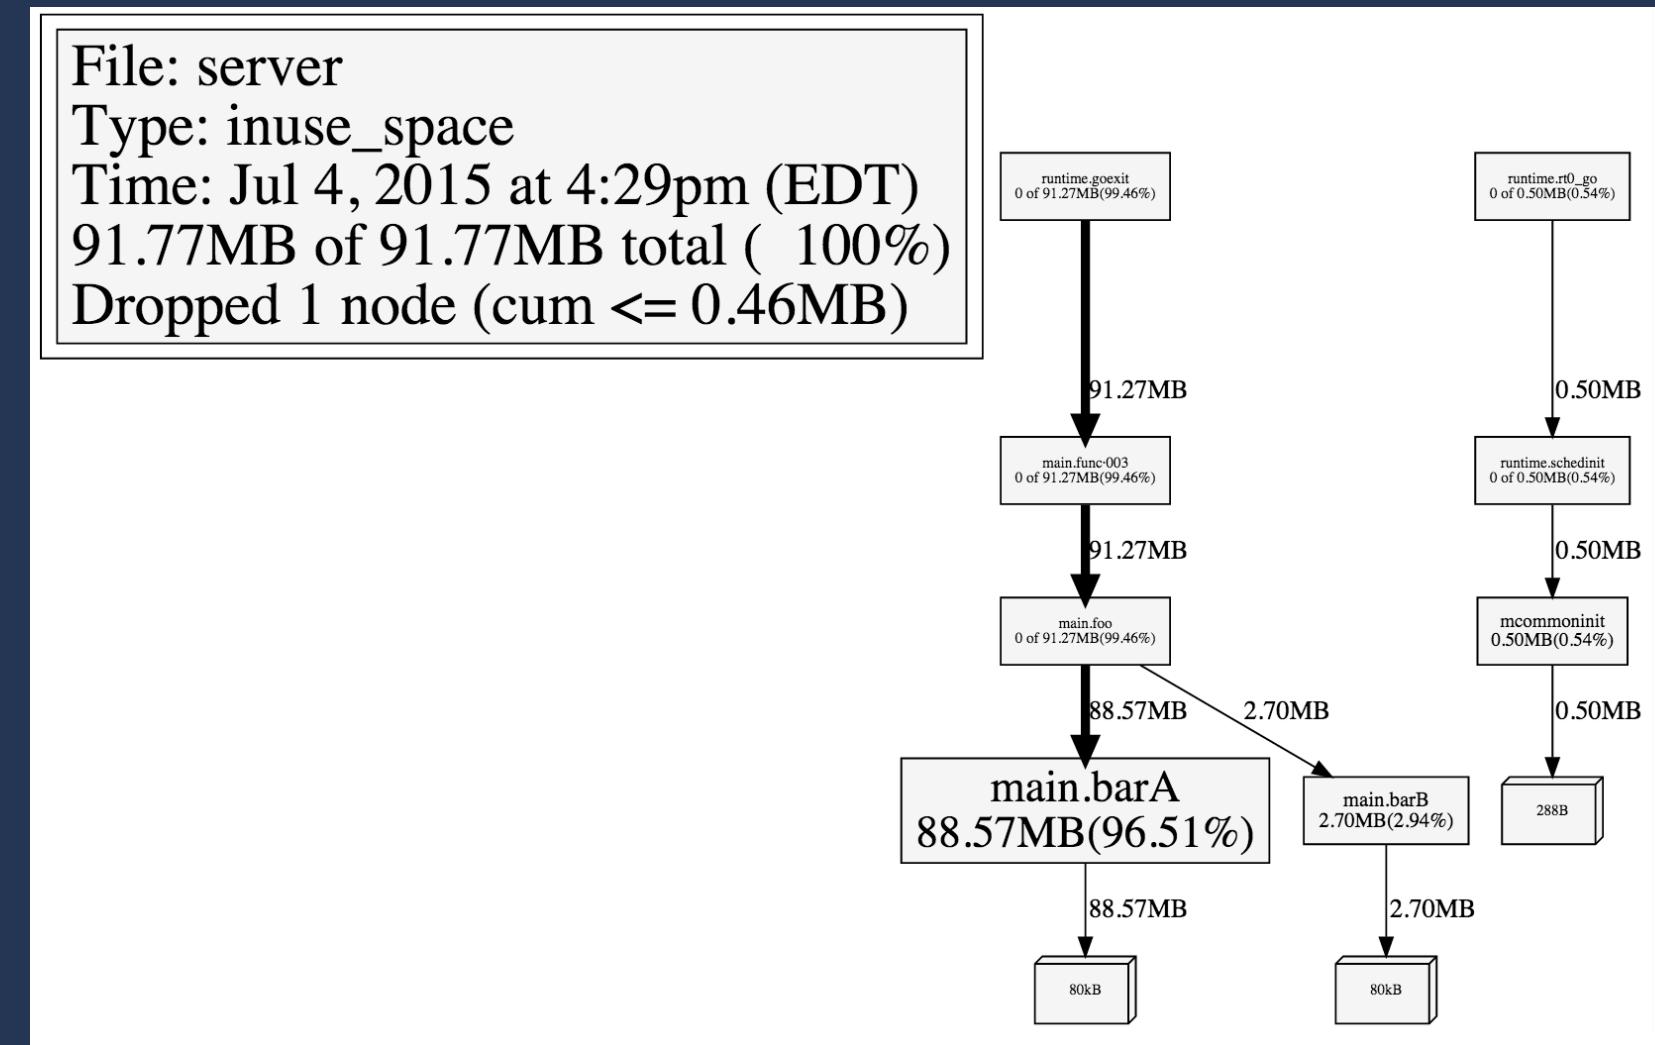
/connections Tells Me: Initech is connected 10 times

```
{  
    "CurrentUserCount":25,  
    "CurrentlyAuthenticatedUsers":  
    [  
        { "Account":"initech", "Name":"Peter Gibbons", ... },  
        ...  
    ]  
}
```

Dev Profiling Told Me: Each connection uses 4MB of memory



Dev PPROF Told Me: Most of that 4MB is SSH internals



**Wolfram Alpha Tells Me:
4MB x 10 = 40MB**

We Contact Initech.

**Timeouts Make Sure That:
Their connections will be closed**

**Prod Profiling Told Me:
This memory will be reclaimed**



uptime:
Preserved.



thanks!

Blake Caldwell

<http://blakecaldwell.net>

<https://twitter.com/blakecaldwell>

Memory Profiler: <http://github.com/fogcreek/profiler>

Credit: Images

- Thinking Man: خالد منتصف - License: CC BY-SA 4.0 (<http://creativecommons.org/licenses/by-sa/4.0>)
- **Race car**: Newlin Keen, Kollins Communications, Inc